

CONTENTS

1. Introduction	1
1.1. Features	1
1.2. System Requirements	2
1.3. Using this Manual	2
2. Getting Started	3
3. Compiler Structure	5
4. Operating Details	7
5. Specific Features	13
5.1. ANSI C Standard Compatibility	13
5.2. Type Checking	13
5.3. Member Names	13
5.4. Unsigned Types	14
5.5. Arithmetic Operations	14
5.6. Structure Operations	15
5.7. Enumerated Types	16
5.8. Initialization Syntax	16
5.9. Function Prototypes	17
5.10. Void and Pointer to Void	18
5.11. Type qualifiers	19
5.12.	20
5.13. Pragma Directives	20
6. Machine Dependencies	21
6.1. Predefined Macros	21
7. Error Checking and Reporting	23
8. Standard Libraries	25
8.1. Standard I/O	25
8.2. Compatibility	25
8.3. Libraries for Embedded Systems	25
8.4. Binary I/O	26
8.5. Floating Point Library	27
9. Stylistic Considerations	29
9.1. Member Names	29
9.2. Use of Int	30
9.3. Extern Declarations	30
10. Memory Models	31
11. What Went Wrong	33
12. Z80 Assembler Reference Manual	35
12.1. Introduction	35
12.2. Usage	35
12.3. The Assembly Language	36
12.3.1. Symbols	36
12.3.1.1. Temporary Labels	37
12.3.2. Constants	37
12.3.2.1. Character Constants	38

12.3.2.2. Floating Constants	38
12.3.2.3. Opcode Constants	38
12.3.3. Expressions	38
12.3.3.1. Operators	38
12.3.3.2. Relocatability	39
12.3.4. Pseudo-ops	40
12.3.4.1. DEFB, DB	40
12.3.4.2. DEFF	41
12.3.4.3. DEFW	41
12.3.4.4. DEFS	41
12.3.4.5. EQU	41
12.3.4.6. DEFL	41
12.3.4.7. DEFM	42
12.3.4.8. END	42
12.3.4.9. COND, IF, ELSE, ENDC	42
12.3.4.10. ELSE	42
12.3.4.11. ENDC	42
12.3.4.12. ENDM	43
12.3.4.13. PSECT	43
12.3.4.14. GLOBAL	43
12.3.4.15. ORG	44
12.3.4.16. MACRO	44
12.3.4.17. LOCAL	45
12.3.4.18. REPT	46
12.3.5. IRP and IRPC	47
12.3.6. Extended Condition Codes	48
12.4. Assembler Directives	48
12.5. Diagnostics	49
12.6. Z80/Z180/64180 Instruction Set	50
13. Linker Reference Manual	69
13.1. Relocation and Psects	69
13.1.1. Program Sections	69
13.1.2. Local Psects and the Large Model	70
13.2. Global Symbols	70
13.3. Operation	71
13.4. Examples	74
13.5. Invoking the Linker	74
14. Librarian	75
14.1. The Library Format	75
14.2. Using	75
14.3. Examples	76
14.4. Supplying Arguments	77
14.5. Listing Format	77
14.6. Ordering of Libraries	77
14.7. Error Messages	78
15. Objtohex	79
16. Cref	83
APPENDIX 1 Error Messages	85
APPENDIX 2 Standard Library Functions	99
INDEX	165

HI-TECH C COMPILER

User's Manual

March 1989

1. Introduction

The HI-TECH C Compiler is a set of software which translates programs written in the C language to executable machine code programs. Versions are available which compile programs for operation under the host operating system, or which produce programs for execution in embedded systems without an operating system.

1.1. Features

Some of HI-TECH C's features are:

A single command will compile, assemble and link entire programs.

The compiler performs strong type checking and issues warnings about various constructs which may represent programming errors.

The generated code is extremely small and fast in execution.

A full run-time library is provided implementing all standard C input/output and other functions.

The source code for all run-time routines is provided.

A powerful general purpose macro assembler is included.

Programs may be generated to execute under the host operating system, or customized for installation in ROM.

PC-DOS/MS-DOS
CP/M-86
Concurrent DOS
Atari ST
Xenix
Unix
CP/M-80

Table 1. Supported Hosts

1.2. System Requirements

The HI-TECH C Compilers operate under the operating systems listed in table 1. Ensure that the version of the compiler you have matches the system you have. Note that in general you must have a hard disk or two floppy disks on your system (it is possible to use one floppy disk of 800K or more). A hard disk is strongly recommended. Note that the CP/M-80 native compiler does not have all the features described in this manual, as it has not been upgraded past V3.09 due to memory limitations. The Z80 cross compiler does support all of the features described here and can be used to generate programs to execute under CP/M-80.

1.3. Using this Manual

The documentation supplied with the HI-TECH C compiler comprises two separate manuals within the one binder. The manual you are reading now covers all versions of the compiler (reflecting the portable nature of the compiler). A separate manual covers machine dependent aspects of your compiler, e.g. installation.

This manual assumes you are familiar with the C language already. If you are not, you should have at least one reference book covering C, of which a large number are available from most computer bookstores, e.g. "A Book on C" by Kelley and Pohl. Other suitable texts are "Programming in ANSI C" by S. Kochan and "The C Programming Language", by Kernighan and Ritchie. You should read the "Getting Started" chapter in this manual, and the "Installation" chapter in the machine-specific manual. This will provide you with sufficient information to work through the introductory examples in the C reference you are using.

Once you have a basic grasp of the C language, the remainder of this manual will provide you with information to enable you to explore the more advanced aspects of C.

Most of the manual covers all implementations of the HI-TECH C compiler. A separate manual is provided for the macro assembler for your particular machine, and other machine-dependent information.

2. Getting Started

If using the compiler on a hard disk system you will need to install the compiler before using it. See the "Installation" chapter for more details. If using a floppy disk based system, in general you should have a copy of the distribution disk #1 in drive A: and maintain your files on a disk in the B: drive. Again see the "Installation" chapter for more information.

```
main()
{
    printf("Hello, world\n");
}
```

Fig. 1. Sample C Program

Before compiling your program it must be contained in a file with an extension (or file type, i.e. that part of the name after the '.') of .C. For example you may type the program shown in fig. 1 into a file called HELLO.C. You will need a text editor to do this. Generally any text editor that can create a straight ASCII file (i.e. not a "word processor" type file) will be suitable. If using editors such as Wordstar, you should use the "non-document mode". Once you have the program in such a file all that is required to compile it is to issue the C command, e.g. to compile HELLO.C simply type the command

```
C -V HELLO.C
```

Cross compilers (i.e. compilers that operate on one system but produce code for a separate target system) will have a compiler driver named slightly differently, e.g. the 68HC11 cross compiler driver is called C68.

If you are using a floppy disk based system (or a CP/M system) it may be necessary to specify where to find the C command, e.g. if the C command is on a disk in drive A: and you are working on B:, type the command

```
A:C -V HELLO.C
```

The compiler will issue a sign on message, then proceed to execute the various passes of the compiler in sequence to compile the program. If you are using a floppy disk based system where the compiler will not fit on a single disk you will be prompted to change disks whenever the compiler cannot find a pass. In this case you should insert a copy of the next distribution disk in drive A: and press RETURN.

As each pass of the compiler is about to be executed, a command line to that pass will be displayed on the screen.

This is because the -V option has been used. This stands for Verbose, and had it not been given the compilation would have been silent except for the sign on message. Error messages can be redirected to a file by using the standard output redirection notation, e.g. > somefile.

After completion of compilation, the compiler will exit to command level. You will notice that several temporary files created during compilation will have been deleted, and all that will be left on the disk (apart from the original source file HELLO.C) will be an executable file. The name of this executable file will be HELLO.EXE for MS-DOS, HELLO.PRG for the Atari ST, HELLO.COM for CP/M-80 and HELLO.CMD for CP/M-86. For cross compilers it will be called HELLO.HEX or HELLO.BIN depending on the default output format for the particular compiler. To execute this program, simply type

HELLO

and you should be rewarded with the message "Hello, world!" on your screen. If you are using a cross compiler you will need to put the program into EPROM or download to the target system to execute it. Cross compilers do not produce programs executable on the host system.

There are other options that may be used with the C command, but you will not need to use them unless you wish to do so. If you are new to the C language it will be advisable to enter and compile a few simple programs (e.g. drawn from one of the C reference texts mentioned above) before exploring other capabilities of the HI-TECH C compiler.

There is one exception to the above; if you compile a program which uses floating point arithmetic (i.e. real numbers) you MUST specify to the compiler that the floating point library should be searched. This is done with a -LF option at the END of the command line, e.g.

C -V FLOAT.C -LF

3. Compiler Structure

The compiler is made up of several passes; each pass is implemented as a separate program. Note that it is not necessary for the user to invoke each pass individually, as the C command runs each pass automatically. Note that the machine dependent passes are named differently for each processor, for example those with 86 in their name are for the 8086 and those with 68K in their name are for the 68000.

The passes are:

CPP The pre-processor - handles macros and conditional compilation

P1 The syntax and semantic analysis pass. This writes intermediate code for the code generator to read.

CGEN, CG86 etc.

The code generator - produces assembler code.

OPTIM, OPT86 etc.

The code improver - may optionally be omitted, reducing compilation time at a cost of larger, slower code produced.

ZAS, AS86 etc.

The assembler - in fact a general purpose macro assembler.

LINK

The link editor - links object files with libraries.

OBJTOHEX

This utility converts the output of LINK into the appropriate executable file format (e.g. .EXE or .PRG or .HEX).

The passes are invoked in the order given. Each pass reads a file and writes a file for its successor to read. Each intermediate file has a particular format; CPP produces C code without the macro definitions and with uses of macros expanded; P1 writes a file containing a program in an intermediate code; CGEN translates this to assembly code; AS produces object code, a binary format containing code bytes along with relocation and symbol information. LINK accepts object files and libraries of object files and writes another object file; this may be in absolute form or it may preserve relocation information and be input to another LINK command.

There are also other utility programs:

LIBR

Creates and maintains libraries of object modules

CREF

Produces cross-reference listings of C or assembler programs.

4. Operating Details

HI-TECH C was designed for ease of use; a single command will compile, assemble and link a C program. The syntax of the C command is as follows:

```
C [ options ] files [ libraries ]
```

The options are zero or more options, each consisting of a dash ('-'), a single key letter, and possibly an argument following the key letter with no intervening space. The files are one or more C source files, assembler source files, or object files. Libraries may be zero or more library names, or the abbreviated form -lname which will be expanded to the library name libname.lib.

The C command will, as determined by the specified options, compile any C source files given, assemble them into object code unless requested otherwise, assemble any assembler source files specified, then link the results of the assemblies with any object or library files specified.

If the C command is invoked without arguments, then it will prompt for a command line to be entered. This command line may be extended by typing a backslash ('\') on the end of the line. Another line will then be requested. If the standard input of the command is from a file (e.g. by typing C < afile) then the command lines will be read from that file. Within the file more than one line may be given if each line but the last ends with a backslash. Note that this mechanism does not work in MS-DOS batch file, i.e. the command file for the C command must be a separate file. MS-DOS has no mechanism for providing long command lines or standard input from inside a batch file.

The options recognized by the C command are as follows:

- S Leave the results of compilation of any C files as assembler output. C source code will be interspersed as comments in the assembler code.
- C Leave the results of all compiles and assemblies as object files; do not invoke the linker. This allows the linker to be invoked separately, or via the C command at a later stage.
- CR Produce a cross reference listing. -CR on its own will leave the raw cross-reference information in a temporary file, allowing the user to run CREF explicitly, while supplying a file name, e.g. -CRFRED.CRF will cause CREF to be invoked to process the raw information into the specified file, in this case FRED.CRF.
- CPM
For the Z80 cross compiler only, produce CP/M-80 COM files. Unless the -CPM option is given, the Z80 cross

compiler uses the ROM runtime startoff module and produces hex or binary images. If the -CPM option is given, CP/M-80 runtime startoff code is linked used and a CP/M-80 COM file is produced.

-O Invoke the optimizer on all compiled code; also requests the assembler to perform jump optimization.

-OOUTFILE

Specify a name for the executable file to be created. By default the name for the executable file is derived from the name of the first source or object file specified to the compiler. This option allows the default to be overridden. If no dot ('.') appears in the given file name, an extension appropriate for the particular operating system will be added, e.g. -OFRED will generate a file FRED.EXE on MS-DOS or FRED.COM on CP/M-86. For cross compilers this also provides a means for specifying the output format, e.g. specifying an output file PROG.BIN will make the compiler generate a binary file, while specifying PROG.HEX will make it generate a hexadecimal file.

-V Verbose: each step of the compilation will be echoed as it is executed.

-I Specify an additional filename prefix to use in searching for #include files. For CP/M the default prefix is 0:A: (user number 0, disk drive A). For MS-DOS the default prefix is A:\HITECH\. Under Unix and Xenix the default prefix is /usr/hitech/include/. Note that on MS-DOS a trailing backslash must be appended to any directory name given as an argument to -I; e.g. -I\FRED\ not -I\FRED. Under Unix a trailing slash should be added.

-D Define a symbol to the preprocessor: e.g. -DCPM will define the symbol CPM as though via #define CPM 1.

-U Undefine a pre-defined symbol. The complement of -D.

-F Request the linker to produce a symbol file, for use with the debugger.

-R For the Z80 CP/M compiler only this option will link in code to perform command line I/O redirection and wild card expansion in file names. See the description of _getargs() in appendix 5 for details of the syntax of the redirections.

-X Strip local symbols from any files compiled, assembled or linked. Only global symbols will remain.

-M Request the linker to produce a link map.

-A This option, for the Z80 only, will cause the compiler to produce an executable program that will, on execution, self-relocate itself to the top of the TPA

(Transient Program Area). This allows the writing of programs which may execute other programs under themselves. Note that a program compiled in such a manner will not automatically reset the bdos address at location 6 in order to protect itself. This must be done by the program itself.

For cross compilers this provides a way of specifying to the linker the addresses that the compiled program is to be linked at. The format of the option is -AROMADR, RAMADR, RAMSIZE. ROMADR is the address of the ROM in the system, and is where the executable code and initialized data will be placed. RAMADR is the start address of RAM, and is where the bss psect will be placed, i.e. uninitialized data. RAMSIZE is the size of RAM available to the program, and is used to set the top of the stack.

For the 6801/6301/68HC11 compiler, the -A option takes a fourth value which is the address of a four byte direct page area called ctemp which the compiled code uses as a scratch pad. If the ctemp address is omitted from the -A option, it defaults to address 0. Normally this will be acceptable, however some 6801 variants (like the 6303) have memory mapped I/O ports at address 0 and start their direct page RAM at address \$80.

For the large memory model of the 8051 compiler, the -A option takes the form -AROMADR, INTRAM, EXTRAM, EXTSIZE. ROMADR is the address of ROM in the system. INTRAM is the start address of internal RAM, and is where the rbss psect will be placed. The 8051 internal stack will start after the end of the rbss psect. EXTRAM is the start address of external RAM, and is where the bss psect will be placed. EXTSIZE is the size of external RAM available to the program, and is used to set the top of the external stack.

- B For compilers which support more than one "memory model", this option is used to select which memory model code is to be generated for. The format of this option is -Bx where x is one or more letters specifying which memory model to use. For the 8086, this option is used to select which one of five memory models (Tiny, Small, Medium, Compact or Large) is to be used.

For the 8051 compiler this this option is used to select which one of the three memory models (Small, Medium or Large) is to be used. For the 8051 compiler only, this option can also be used to select static allocation of auto variables by appending an A to the end of the -B option. For example, -Bsa would select small model with static allocation of all variables, while -Bm would select medium model with auto variables dynamically allocated on the stack.

- E By default the 8086 compiler will initialize the executable file header to request a 64K data segment at run time. This may be overridden by the -E option. It takes an argument (usually in hexadecimal representation) which is the number of BYTES (not paragraphs) to be allocated to the program at run time. For example -E0ffff0h will request a megabyte. Since this much will not be available, the operating system will allocate as much as it can.

- W This options sets the warning level, i.e. it determines how picky the compiler is about legal but dubious type conversions etc. -W0 will allow all warning messages (default), -W1 will suppress the message "Func() declared implicit int". -W3 is recommended for compiling code originally written with other, less strict, compilers. -W9 will suppress all warning messages.

- H This option generates a symbol file for use with debuggers. The format of the symbol file is described elsewhere. The default name of the symbol file is l.sym. An alternate name may be specified with the option, e.g. -Hsymfile.abc.

- G Like -H, -G also generates a symbol file, but one that contains line and file number information for a source level debugger. Like -H a file name may be specified. When used in conjunction with -O, only partial optimization will be performed to avoid confusing the debugger.

- P Execution profiling is available on native compilers running under DOS, CP/M-86 and on the Atari ST. This option generates code to turn on execution profiling when the program is run. A -H option should also be specified to provide a symbol table for the profiler EPROF.

- Z For version 5.xx compilers only, the -Z option is used to select global optimization of the code generated. For the 8086 and 6801/6301/68HC11 compilers the only valid -Z option is -Zg. For the 8051 compiler, the valid -Z options are -Zg which invokes global optimization, -Zs which optimizes for space, and -Zf which optimizes for speed. The s and f options may be used with the g option, thus the options -Zgf and -Zgs are valid. Speed and space optimization are mutually exclusive, i.e. the s and f options cannot be used together.

- 1 For the 8086 compiler only, request the generation of code which takes advantage of the extra instructions of the 80186 processor. A program compiled with -1 will not execute on an 8086 or 8088 processor. For the 68000 compiler, generate instructions for the 68010 processor.

- 2 Like -1 but for the 80286 and 68020.
- 11 For the 6801/HC11 compiler, this option will request generation of instructions specific to the 68HC11 processor.
- 6301
For the 6801/HC11 compiler, this option will request generation of instructions specific to the 6301/6303 processors.

Some examples of the use of the C command are:

```
c prog.c  
  
c -mlink.map prog.c x.obj -lx  
  
c -S prog.c  
  
c -O -C -CRprog.crf prog.c prog2.c  
  
c -v -Oxfile.exe afile.obj afile.c -lf
```

Upper and lower case has been used in the above examples for emphasis: the compiler does not distinguish between cases, although arguments specifying names, e.g. -D, are inherently case sensitive.

Taking the above examples in order; the first compiles and links the C source file prog.c with the standard C library. The second example will compile the file prog.c and link it with the object file x.obj and the library libx.lib; a link map will be written to the file link.map.

The third example compiles the file prog.c, leaving the assembler output in a file prog.as. It does not assemble this file or invoke the linker. The next example compiles both prog.c and prog2.c, invoking the optimizer on both files, but does not perform any linking. A cross reference listing will be left in the file prog.crf.

The last example pertains to the 8086 version of the compiler. It runs the compiler with the verbose option, and will cause afile.c to be compiled without optimization to object code, yielding afile.obj, then afile.obj and afile.obj will be linked together with the floating point library (from the -LF option) and the standard library to yield the executable program xfile.exe (assuming this is performed on an MS-DOS system). One would expect this program to use floating point, if it did not then the -LF option would have been required.

If more than one C or assembler source file is given to the C command, the name of each file will be printed on the console as it is processed. If any fatal errors occur during the compilation or assembly of source files, further source files will be processed, but the linker will not be

invoked.

Other commands which may be issued by the user, rather than automatically by the C command, are:

ZAS The Z80 assembler.

AS86
The 8086 assembler.

LINK
The linker

LIBR
The library maintainer

OBJTOHEX
Object to hex converter

CREF
Cross reference generator.

In general, these commands accept the same type of command line as the C command, i.e. zero or more options (indicated by a leading '-') followed by one or more file arguments. If the linker or the librarian is invoked with no arguments, it will prompt for a command line. This allows command lines of more than 128 bytes to be entered. Input may also be taken from a file by using the redirection capabilities (see `_getargs()` in the library function listing). See the discussion above of the C command. These commands are described in further detail in their respective manuals.

5. Specific Features

The HI-TECH C compiler has a number of features, which while largely compatible with other C compilers, contribute to more reliable programming methods.

5.1. ANSI C Standard Compatibility

At the time of writing the Draft ANSI Standard for the C Language was at an advanced stage, though not yet an official standard. Accordingly it is not possible to claim compliance with that standard, however HI-TECH C includes the majority of the new and altered features in the draft ANSI standard. Thus it is in the sense that most people understand it "ANSI compatible".

5.2. Type Checking

Previous C compilers have adopted a lax approach to type checking. This is typified by the Unix C compiler, which allows almost arbitrary mixing of types in expressions. The HI-TECH C compiler performs much more strict type checking, although in most cases only warning messages are issued, allowing compilation to proceed if the user knows that the errors are harmless. This would occur, for example, when an integer value was assigned to a pointer variable. The generated code would almost certainly be what the user intended, however if in fact it represented an error in the source code, the user is prompted to check and correct it where necessary.

5.3. Member Names

In early C compilers member names in different structures were required to be distinct except under certain circumstances. HI-TECH C, like most recent implementations of C, allows member names in different structures and unions to overlap. A member name is recognized only in the context of an expression whose type is that of the structure in which the member is defined. In practice this means that a member name will be recognized only to the right of a '.' or '->' operator, where the expression to the left of the operator is of type structure or pointer to structure the same as that in which the member name was declared. This not only allows structure names to be re-used without conflict in more than one structure, it permits strict checking of the usage of members; a common error with other C compilers is the use of a member name with a structure pointer of the wrong type, or worse with a variable which is a pointer to a simple type.

There is however an escape from this, where the user desires to use as a structure pointer something which is not declared as such. This is via the use of a typecast. For example, suppose it is desired to access a memory-mapped i/o device, consisting of several registers. The declarations

and use may look something like the code fragment in fig. 2.

```

struct io_dev
{
    short    io_status;    /* status */
    char     io_rxddata;    /* rx data */
    char     io_txddata;    /* tx data */
};

#define RXRDY    01        /* rx ready */
#define TXRDY    02        /* tx ready */

/* define the (absolute) device address */

#define DEVICE    ((struct io_dev *)0xFF00)

send_byte(c)
char c;
{
    /* wait till transmitter ready */
    while(!(DEVICE->io_status & TXRDY))
        continue;
    /* send the data byte */
    DEVICE->io_txddata = c;
}

```

Fig. 2. Use of Typecast on an Absolute Address

In this example, the device in question has a 16 bit status port, and two 8 bit data ports. The address of the device (i.e. the address of its status port) is given as (hex)0FF00. This address is typecast to the required structure pointer type to enable use of the structure member names. The code generated by this will use absolute memory references to access the device, as required.

Some examples of right and wrong usage of member names are shown in fig. 3.

5.4. Unsigned Types

HI-TECH C implements unsigned versions of all integral types; i.e. unsigned char, short, int and long. If an unsigned quantity is shifted right, the shift will be performed as a logical shift, i.e. bringing zeros into the rightmost bits. Similarly right shifts of a signed quantity will sign extend the rightmost bits.

5.5. Arithmetic Operations

On machines where arithmetic operations may be performed more efficiently in lengths shorter than int, operands shorter than int will not be extended to int length unless necessary.

For example, if two characters are added and the result stored into another character, it is only necessary to


```
struct fred
{
    char    a;
    int     b;
}    s1, * s2;

struct bill
{
    float    c;
    long     b;
}    x1, * x2;

main()
{
    /* wrong - c is not a member of fred */
    s1.c = 2;

    /* correct */
    s1.a = 2;

    /* wrong - s2 is a pointer */
    s2.a = 2;

    /* correct */
    x2->b = 24L;

    /* right, but note type conversion
       from long to int */
    s2->b = x2->b;
}
```

Fig. 3. Examples of Member Usage

perform arithmetic in 8 bits, since any overflow into the top 8 bits will be lost. However, if the sum of two characters is stored into an int, the addition should be done in 16 bits to ensure the correct result.

In accordance with the draft ANSI standard, operations on float rather than double quantities will be performed in the shorter precision rather than being converted to double precision then back again.

5.6. Structure Operations

HI-TECH C implements structure assignments, structure arguments and structure-valued functions in their full generality. The example in fig. 4 is a function returning a structure. Some legal (and illegal) uses of the function are also shown.

```

struct bill
{
    char    a;
    int     b;
}
afunc()
{
    struct bill    x;

    return x;
}

main()
{
    struct bill    a;

    a = afunc();          /* ok */
    pf("%d", afunc().a);  /* ok */

    /* illegal, afunc() cannot be assigned
       to, therefore neither can
       afunc().a */
    afunc().a = 1;

    /* illegal, same reason */
    afunc().a++;
}

```

Fig. 4. Example of a Function Returning a Structure

5.7. Enumerated Types

HI-TECH C supports enumerated types; these provide a structured way of defining named constants.

The uses of enumerated types are more restricted than that allowed by the Unix C compiler, yet more flexible than permitted by LINT. In particular, an expression of enumerated type may be used to dimension arrays, as an array index or as the operand of a switch statement. Arithmetic may be performed on enumerated types, and enumerated type expressions may be compared, both for equality and with the relation operators. An example of the use of an enumerated type is given in fig. 5.

5.8. Initialization Syntax

Kernighan and Ritchie in "The C Programming Language" state that pairs of braces may be omitted from an initializer in certain contexts; the draft ANSI standard provides that a conforming C program must either include all braces in an initializer, or leave them all out. HI-TECH C allows any pairs of braces to be omitted providing that the front end of the compiler can determine the size of any arrays being initialized, and providing that there is no ambiguity as to which braces have been omitted. To avoid ambiguity if any pairs of braces are present then any braces which would

```

/* a represents 0, b -> 1 */
enum fred { a, b, c = 4 };

main()
{
    enum fred      x, y, z;

    x = z;
    if(x < z)
        func();
    x = (enum fred)3;
    switch(z) {
    case a:
    case b:
    default:
    }
}

```

Fig. 5. Use of an Enumerated Type

enclose those braces must also be present. The compiler will complain ("initialization syntax") if any ambiguity is present.

5.9. Function Prototypes

A new feature of C included in the proposed ANSI for C, known as "function prototypes", provides C with an argument checking facility, i.e. it allows the compiler to check at compile time that actual arguments supplied to a function invocation are consistent with the formal parameters expected by the function. The feature allows the programmer to include in a function declaration (either an external declaration or an actual definition) the types of the parameters to that function. For example, the code fragment shown in fig. 6 shows two function prototypes.

```

void fred(int, long, char *);

char *
bill(int a, short b, ...)
{
    return a;
}

```

Fig. 6. Function Prototypes

The first prototype is an external declaration of the function `fred()`, which accepts one integer argument, one long argument, and one argument which is a pointer to `char`. Any usage of `fred()` while the prototype declaration is in scope will cause the actual parameters to be checked for number and type against the prototype, e.g. if only two arguments were supplied or an integral value was supplied for the third argument the compiler would report an error.

In the second example, the function `bill()` expects two or more arguments. The first and second will be converted to `int` and `short` respectively, while the remainder (if present) may be of any type. The ellipsis symbol (`...`) indicates to the compiler that zero or more arguments of any type may follow the other arguments. The ellipsis symbol must be last in the argument list, and may not appear as the only argument in a prototype.

All prototypes for a function must agree exactly, however it is legal for a definition of a function in the old style, i.e. with just the parameter names inside the parentheses, to follow a prototype declaration provided the number and type of the arguments agree. In this case it is essential that the function definition is in scope of the prototype declaration.

Access to unspecified arguments (i.e. arguments supplied where an ellipsis appeared in a prototype) must be via the macros defined in the header file `<stdarg.h>`. This defines the macros `va_start`, `va_arg` and `va_end`. See `va_start` in the library function listing for more information.

NOTE that it is a grave error to use a function which has an associated prototype unless that prototype is in scope, i.e. the prototype MUST be declared (possibly in a header file) before the function is invoked. Failure to comply with this rule may result in strange behaviour of the program. HI-TECH C will issue a warning message ("func() declared implicit int") whenever a function is called without an explicit declaration. It is good practice to declare all functions and global variables in one or more header files which are included wherever the functions are defined or referenced.

5.10. Void and Pointer to Void

The void type may be used to indicate to the compiler that a function does not return a value. Any usage of the return value from a void function will be flagged as an error.

The type `void *`, i.e. pointer to void, may be used as a "universal" pointer type. This is intended to assist in the writing of general purpose storage allocators and the like, where a pointer is returned which may be assigned to another variable of some other pointer type. The compiler permits without typecasting and without reporting an error the conversion of `void *` to any other pointer type and vice versa. The programmer is advised to use this facility carefully and ensure that any `void *` value is usable as a pointer to any other type, e.g. the alignment of any such pointer should be suitable for storage of any object.

5.11. Type qualifiers

The ANSI C standard introduced the concept of type qualifiers to C; these are keywords that qualify the type to which they are applied. The type qualifiers defined by ANSI C are `const` and `volatile`. HI-TECH C also implements several other type qualifiers. The extra qualifiers include:

```
far
near
interrupt
fast interrupt
port
```

Not all versions of the compilers implement all of the extra qualifiers. See the machine dependent section for further information.

When constructing declarations using type qualifiers, it is very easy to be confused as to the exact semantics of the declaration. A couple of rules-of-thumb will make this easier. Firstly, where a type qualifier appears at the left of a declaration it may appear with any storage class specifier and the basic type in any order, e.g.

```
static void interrupt func();
```

is semantically the same as

```
interrupt static void func();
```

Where a qualifier appears in this context, it applies to the basic type of the declaration. Where a qualifier appears to the right of one or more '*' (star) pointer modifiers, then you should read the declaration from right to left, e.g.

```
char * far fred;
```

should be read as "fred is a far pointer to char". This means that fred is qualified by far, not the char to which it points. On the other hand,

```
char far * bill;
```

should be read as "bill is a pointer to a far char", i.e. the char to which bill points is located in the far address space. In the context of the 8086 compiler this will mean that bill is a 32 bit pointer while fred is a 16 bit pointer. You will hear bill referred to as a "far pointer", however the terminology "pointer to far" is preferred.

5.12.

There are two methods provided for in-line assembler code in C programs. The first allows several lines of assembler anywhere in a program. This is via the `#asm` and `#endasm` preprocessor directives. Any lines between these two directives will be copied straight through to the assembler file produced by the compiler. Alternatively you can use the `asm("string");` construct anywhere a C statement is expected. The string will be copied through to the assembler file. Care should be taken with using in-line assembler since it may interact with compiler generated code.

5.13. Pragma Directives

The draft ANSI C standard provides for a `#pragma` preprocessor directive that allows compiler implementations to control various aspects of the compilation process. Currently HI-TECH C only supports one pragma, the `pack` directive. This allows control over the manner in which members are allocated inside a structure. By default, some of the compilers (especially the 8086 and 68000 compilers) will align structure members onto even boundaries to optimize machine accesses. It is sometimes desired to override this to achieve a particular layout inside a structure. The `pack` pragma allows specification of a maximum packing factor. For example, `#pragma pack1(1)` will instruct the compiler that no additional padding be inserted between structure members, i.e. that all members should be aligned on boundaries divisible by 1. Similarly `#pragma pack(2)` will allow alignment on boundaries divisible by 2. In no case will use of the `pack` pragma force a greater alignment than would have been used for that data type anyway.

More than one `pack` pragma may be used in a program. Any use will remain in force until changed by another `pack` or until the end of the file. Do not use a `pack` pragma before include files such as `<stdio.h>` as this will cause incorrect declarations of run-time library data structures.

6. Machine Dependencies

HI-TECH C eliminates many of the machine dependent aspects of C, since it uniformly implements such features as unsigned char. There are however certain areas where machine dependencies are inherent in the C language; programmers should be aware of these and take them into account when writing portable code.

The most obvious of these machine dependencies is the varying size of C types; on some machines an int will be 16 bits, on others it may be 32 bits. HI-TECH C conforms to the following rules, which represent common practice in most C compilers.

char	is at least 8 bits
short	is at least 16 bits
long	is at least 32 bits
int	is the same as either short or long
float	is at least 32 bits
double	is at least as wide as float

Because of the variable width of an int, it is recommended that short or long be used wherever possible in preference to int. The exception to this is where a quantity is required to correspond to the natural word size of the machine.

Another area of machine differences is that of byte ordering; the ordering of bytes in a short or long can vary significantly between machines. There is no easy approach to this problem other than to avoid code which depends on a particular ordering. In particular you should avoid writing out whole structures to a file (via fwrite()) unless the file is only to be read by the same program then deleted. Different compilers use different amounts of padding between structure members, though this can be modified via the #pragma pack(n) construct.

6.1. Predefined Macros

One technique through which machine unavoidable machine dependencies may be managed is the predefined macros provided by each compiler to identify the target processor and operating system (if any). These are defined by the compiler driver and may be tested with conditional compilation preprocessor directives.

The macros defined by various compilers are listed in table 2. These can be used as shown in the example in fig .

M a c_r_o	D_e_f_i_n_e_d_f_o_r
i8051	8051 processor family
i8086	8086 processor family
i8096	8096 processor family
z80	Z80 processor and derivatives
m68000	68000 processor family
m6800	6801, 68HC11 and 6301 processors
m6809	6809 processor
DOS	MS-DOS and PC-DOS
CPM	CP/M-80 and CP/M-86
TOS	Atari ST

Table 2. Predefined Macros

```

#if      DOS
char *   filename = "c:file";
#endif   /* DOS */
#if      CPM
char *   filename = "0:B:afile";
#endif   /* CPM */

```

Use this page for notes

7. Error Checking and Reporting

Errors may be reported by any pass of the compiler, however in practice the assembler and optimizer will not encounter any errors in the generated code. The types of errors produced by each pass are summarized below. In general any error will be identified by the name of the source file in which it was encountered, and the line number at which it was detected. P1 will also nominate the name of the function inside which the error was detected.

Errors may be redirected to a file with the usual syntax, i.e. a 'greater than' symbol ('>') followed by the name of a file into which the errors should be written. The file name may of course be a device name, e.g. LST: for CP/M or PRN for MS-DOS.

CPP will report errors relating to macro definitions and expansions, as well as conditional compilation.

P1 is the pass which reports most errors; it performs syntax and semantic checking of the input, and will report both fatal and warning errors when encountered. Syntax errors will normally be expressed as "symbol expected" or "symbol unexpected". Semantic errors may relate to undeclared, redeclared or misdeclared variables. P1 will also report variable definitions which are unused or unreferenced. These errors are warnings, as are most type checking errors.

When P1 encounters errors it will list the source line containing the error on the screen, and underneath display the error message and an up arrow pointing to the point at which the compiler detected the error. In some cases the actual cause of the error may be earlier in the line or even on previous line.

CGEN may occasionally report errors, usually warnings, and mostly related to unusual combinations of types with constants, for example testing if an unsigned quantity is less than zero. One fatal error produced by CGEN is "can't generate code for this expression" and means that the expression currently being compiled is in some way too complicated to produce code for. This can usually be overcome by rewriting the source code. Such errors are rare and will occur only for unusual constructs.

The linker will report undefined or multiply defined symbols. Note that variable declarations inside a header file which is included in more than one source file must be declared as extern, to avoid multiply defined symbol errors. These symbols must then be defined in one and one only source file.

A comprehensive list of error messages is included in an appendix.

Use this page for notes

8. Standard Libraries

8.1. Standard I/O

C is a language which does not specify the I/O facilities in the language itself; rather all I/O is performed via library routines. In practice this results in I/O handling which is no less convenient than any other language, with the added possibility of customising the I/O for a specific application. For example it is possible to provide a routine to replace the standard `getchar()` (which gets one character from the standard input) with a special `getchar()`. This is particularly useful when writing code which is to run on a special hardware configuration, while maintaining a high level of compatibility with "standard" C I/O.

There is in fact a Standard I/O library (STDIO) which defines a portable set of I/O routines. These are the routines which are normally used by any C application program. These routines, along with other non-I/O library routines, are listed in detail in a subsequent section of the manual.

8.2. Compatibility

The libraries supplied with HI-TECH C are highly compatible with the ANSI libraries, as well as the V7 UNIX libraries, both at the Standard I/O level and the UNIX system call level. The Standard I/O library is complete, and conforms in all respects with the UNIX Standard I/O library. The library routines implementing UNIX system call like functions are as close as possible to those system calls, however there are some UNIX system calls that cannot be simulated on other systems, e.g. the `link()` operation. However the basic low level I/O routines, i.e. `open`, `close`, `read`, `write` and `lseek` are identical to the UNIX equivalents. This means that many programs written to run on UNIX, even if they do not use Standard I/O, will run with little modification when compiled with HI-TECH C.

8.3. Libraries for Embedded Systems

The cross compilers, designed to produce code for target systems without operating systems, are supplied with libraries implementing a subset of the STDIO functions. Not provided are those functions dealing with files. Included are `printf()`, `scanf()` etc. These operate by calling two low level functions `putch` and `getch()` which typically send and receive characters via a serial port. Where the compiler is aimed at a single-chip micro with on-board UART this is used. The source code for all these functions is supplied enabling the user to modify it to address a different serial port.

8.4. Binary I/O

On some operating systems, notably CP/M, files are treated differently according to whether they contain ASCII (i.e. printable) or binary data. MD-DOS also suffers from this problem, not due to any lack in the operating system itself, but rather because of the hangover from CP/M which results in many programs putting a redundant ctrl-Z at the end of files. Unfortunately there is no way to determine which type of data a file contains (except perhaps by the name or type of the file, and this is not reliable). To overcome this difficulty, there is an extra character which may be included in the MODE string to a `fopen()` call. To open a file for ASCII I/O, the form of the `fopen()` call is:

```
fopen("filename.ext", "r")    /* for reading */
fopen("filename.ext", "w")    /* for writing */
```

To open a file for binary I/O, the character 'b' may be appended to the

```
fopen("filename.ext", "rb")
fopen("filename.ext", "wb")
```

The additional character instructs the `STDIO` library that this file is to be handled in a strict binary fashion. On CP/M or MS-DOS, a file opened in ASCII mode will have the following special character handling performed by the `STDIO` routines:

`newline`
(''\n'') converted to carriage return/newline on output.

`return`
(''\r'') ignored on input

`Ctrl-Z`
interpreted as End-Of-File on input and, for CP/M only, appended when closing the file.

The special actions performed on ASCII files ensure that the file is written in a format compatible with other programs handling text files, while eliminating the requirement for any special handling by the user program - the file appears to the user program as though it were a UNIX-like text file.

None of these special actions are performed on a file opened in binary mode. This is required when handling any kind of binary data, to ensure that spurious bytes are not inserted, and premature EOF's are not seen.

Since the binary mode character is additional to the normal mode character, this usage is quite compatible with UNIX C. When compiled on UNIX, the additional character will be ignored.

A mention here of the term 'stream' is appropriate; stream is used in relation to the STDIO library routines to mean the source or sink of bytes (characters) manipulated by those routines. Thus the FILE pointer supplied as an argument to the STDIO routines may be regarded as a handle on the corresponding stream. A stream may be viewed as a featureless sequence of bytes, originating from or being sent to a device or file or even some other indeterminate source. A FILE pointer should not be confused with the 'file descriptors' used with the low-level I/O functions open(), close(), read() and write(). These form an independent group of I/O functions which perform unbuffered reads and writes to files.

8.5. Floating Point Library

HI-TECH C supports floating point as part of the language, however the Z80 implementation provides single precision only; double floats are permitted but are no different to floats. In addition, the standard library, LIBC.LIB, does not contain any floating point routines. These have been separated out into another library, LIBF.LIB. This means that if this library is not searched no floating point support routines will be linked in, thus avoiding any size penalty for the floating point support if it is not used. This is particularly important for printf and scanf, and thus LIBF.LIB contains versions of printf and scanf that do support floating point formats.

Thus, if floating point is used, a -LF option should be used AFTER the source and/or object files to the C command. E.g.:

```
C -V -O x.c y.c z.obj -LF
```

Use this page for notes

9. Stylistic Considerations

Although it is not the purpose of this manual to set out a coding standard for C, some comments regarding use of some of the features of HI-TECH C may be useful.

9.1. Member Names

Although HI-TECH C allows the same structure or union member name to be used in more than one structure or union, this may not be allowed by another C compiler. To help ensure portability of the code, it is recommended that member names all be distinct, and a useful way of ensuring this is to prefix each member name with one or two letters derived from the name of the structure itself. An example is given in fig. 7.

```
struct tree_node
{
    struct tree_node *    t_left;
    struct tree_node *    t_right;
    short                 t_operator;
};
```

Fig. 7. Member Naming

Because HI-TECH C insists on use of all intermediate names when referencing a member nested inside several structures, some simple macro definitions can serve as a shorthand. An example is given in fig. 8.

```
struct tree_node
{
    short          t_operator;
    union
    {
        struct tree_node * t_un_sub[2];
        char *              t_un_name;
        long                t_un_val;
    } t_un;
};

#define t_left  t_un.t_un_sub[0]
#define t_right t_un.t_un_sub[1]
#define t_name  t_un.t_un_name
#define t_val   t_un.t_un_val
```

Fig. 8. Member Name Shorthand

This enables the variant components of the structure to be referred to by short names, while guaranteeing portability and presenting a clean definition of the structure.

9.2. Use of Int

It is recommended that the type `int` be avoided wherever possible, in preference to the types `short` or `long`. This is because of the variable size of `int`, whereas `short` is commonly 16 bits and `long` 32 bits in most C implementations.

9.3. Extern Declarations

Some compilers permit a non-initialized global variable to be declared in more than one place, with the multiple definitions being resolved by the linker as all defining the same thing. HI-TECH C specifically disallows this, since it may lead to subtle bugs. Instead, global variables may be declared `extern` in as many places as you wish, and must be defined in one and one only place. Typically this will mean declaring global variables in a header file as `extern`, and defining each variable in the file most closely associated with that variable.

This usage will be portable to virtually all other C implementations.

10. Memory Models

With many of the processors supported by the HI-TECH C compilers there are more than one address space accessible to a program. Typically one address space is more economical to access than another, larger address space. Thus it is desirable to be able to tailor a program's use of memory to achieve the greatest economy in addressing (thus minimizing program size and maximizing speed) while allowing access to as much memory as the program requires.

This concept of different address spaces is not catered for by either K&R or ANSI C (except to recognize the possibility of separate address spaces for code and data). Without any extensions to the language itself it is possible to devise more than one memory model for a given processor, selected at compile time. This has the effect of selecting one addressing method for all data and/or code. This permits the model for a particular program to be chosen depending on that program's memory requirements.

In many programs, however, only one or two data structures are large enough to need to be placed in the larger address space. Selection of a "large" memory model for the whole of the program makes the whole program larger and slower just to allow a few large data structures. This can be overcome by allowing individual selection of the address space for each data structure. Unfortunately this entails extensions to the language, never a desirable approach. To minimize the effect of such extensions they should satisfy the following criteria:

1. As far as possible the extensions should be consistent with common practice.
2. The extensions should fit a machine-independent model to maximize portability across processors and operating systems.

These goals have been achieved within HI-TECH C by means of the following model:

Each memory model defines three address spaces each for code and data. These address spaces are known as the near, far and default spaces. Any object qualified by the near keyword will be placed in the near address space, any object qualified by the far keyword shall be placed in the far address space, and all other objects shall be placed in the default address space. The near address space shall be a (possibly improper) subspace of the default address space, while the default address space shall be a (possibly improper) subspace of the far address space. There shall be up to three kinds of pointers corresponding to the three address spaces, each capable of

addressing an object in its own address space or a subspace of that address space.

This implies that the address of an object may be converted to a pointer into a larger address space, e.g. a near object may have its address converted to a pointer to far, but a far object may not be able to be addressed by a pointer to near.

In practice the default address space will usually correspond exactly to either the near or far address spaces. If all three address spaces correspond to the same memory then there is only one memory model possible. This occurs with the 68000 processor. Where the default code and data spaces may each correspond to either the near or far address spaces then there will be a total of four memory models. This is the case with the 8086 processor.

The keywords far and near are supported by all the HI-TECH C compilers, but the exact correspondence of address spaces is determined by the individual characteristics of each processor and the choice of memory model (if there is a choice). However code written using these keywords will be portable providing it obeys the constraints of the model described above.

This model also corresponds well with other implementations using the near and far keywords, although such implementations do not appear to have been designed around a formal, portable model.

11. What Went Wrong

There are numerous error messages that the compiler may produce. Most of these relate to errors in the source code (syntax errors of various kinds and so forth) but some represent limitations, particularly of memory. The two passes most likely to be affected by memory limitations are the code generator and the optimizer. The code generator will issue the message "No room" if it runs out of dynamic memory. This can usually be eliminated by simplifying the expression at the line nominated in the error message. The more complex the expression, the more memory is required to store the tree representing it. Reducing the number of symbols used in the program will also help.

Note that this error is different from the message "Can't generate code for this expression" which indicates that the expression is in some way too difficult for the code generator to handle. This message will be encountered very infrequently, and can be eliminated by changing the expression in some way, e.g. computing an intermediate value into a temporary variable.

The optimizer reads the assembler code for a whole function into memory at one time. Very large functions will not fit, giving the error message "Optim: out of memory in _func" where func is the name of the function responsible. In this case the function should be broken up into smaller functions. This will only occur with functions with several hundred lines of C source code. Good coding practice will normally limit functions to less than 50 lines each.

If a pass exits with the message "Error closing file", or "Write error on file", this usually indicates that there is insufficient room on the current disk.

If you use a wordprocessing editor such as Wordstar, ensure that you use the "non-document" mode or whatever the corresponding mode is. The edited file should not contain any characters with the high bit set, and the line feed at the end of the line must be present. Lines should be not more than 255 characters long.

When using floating point, ensure that you use a -LF flag at the END of the command line, to cause the floating point library to be searched. This will cause floating versions of printf and scanf to be linked in, as well as specific floating point routines.

If the non-floating version of printf is used with a floating format such as %f then it will simply print the letter f.

If the linker gives an "Undefined symbol" message for some symbol which you know nothing about, it is possible that it is a library routine which was not found during the

library search due to incorrect library ordering. In this case you can search the library twice, e.g. for the standard library add a -LC to the end of the C command line, or -LF for the floating library. If you have specified the library by name simply repeat its name.

12. Z80 Assembler Reference Manual

12.1. Introduction

The assembler incorporated in the HI-TECH C compiler system is a full-featured relocating macro assembler accepting Zilog mnemonics. These mnemonics and the syntax of the Z80 assembly language are described in the "Z80 Assembly Language Handbook" published by Zilog and are included at the end of this manual as a reference. The assembler implements certain extensions to the operands allowed, and certain additional pseudo-ops, which are described here. The assembler also accepts the additional opcodes for the Hitachi 64180 and Z180 processors.

12.2. Usage

The assembler is named `zas`, and is invoked as follows:

`ZAS options files ...`

The files are one or more assembler source files which will be assembled, but note that all the files are assembled as one, not as separate files. To assemble separate files, the assembler must be invoked on each file separately. The options are zero or more options from the following list:

- N Ignore arithmetic overflow in expressions. The -N option suppresses the normal check for arithmetic overflow. The assembler follows the "Z80 Assembly Language Handbook" in its treatment of overflow, and in certain instances this can lead to an error where in fact the expression does evaluate to what the user intended. This option may be used to override the overflow checking.
- J Attempt to optimize jumps to branches. The -J option will request the assembler to attempt to assemble jumps and conditional jumps as relative branches where possible. Only those conditional jumps with branch equivalents will be optimized, and jumps will only be optimized to branches where the target is in branch range. Note that the use of this option slows the assembly down, due to the necessity for the assembler to make an additional pass over the input code.
- U Treat undefined symbols as external. The -U option will suppress error messages relating to undefined symbols. Such symbols are treated as externals in any case. The use of this option will not alter the object code generated, but merely serves to suppress the error messages.

-Ofile

Place the object code in file. The default object file name is constructed from the name of the first source file. Any suffix or file type (i.e. anything following the rightmost dot ('.') in the name is stripped, and the suffix .obj appended. Thus the command

```
ZAS file1.as file2.z80
```

will produce an object file called file1.obj. The use of the -O option will override this default convention, allowing the object file to be arbitrarily named. For example:

```
ZAS -ox.obj file1.obj
```

will place the object code in x.obj.

-Llist

Place an assembly listing in the file list, or on standard output if list is null. A listfile may be produced with the -L option. If a file name is supplied to the option, the list file will be created with that name, otherwise the listing will be written to standard output (i.e. the console). List file names such as CON: and LST: are acceptable.

-Wwidth

The listing is to be formatted for a printer of given width. The -W option specifies the width to which the listing is to be formatted. E.g.

```
ZAS -Llst: -W80 x.as
```

will output a listing formatted for an 80 column printer to the list device.

-C This option requests ZAS to produce cross reference information in a file. The file will be called xxx.crf where xxx is the base part of the first source file name. It will then be necessary to run the CREF utility to turn this information into a formatted listing.

12.3. The Assembly Language

As mentioned above, the assembly language accepted by zas is based on the Zilog mnemonics. You should have some reference book such as the "Z80 Assembly Language Handbook". Described below are those areas in which zas differs, or has extensions, compared to the standard Zilog assembly language.

12.3.1. Symbols

The symbols (labels) accepted by the assembler may be of any length, and all characters are significant. The characters used to form a symbol may be chosen from the upper and lower case alphabets, the digits 0-9, and the special

symbols underscore ('_'), dollar ('\$') and question mark ('?'). The first character may not be numeric. Upper and lower case are distinct. The following are all legal and distinct symbols.

```
An_identifier
an_identifier
an_identifier1
$$$
?$_123455
```

Note that the symbol \$ is special (representing the current location) and may not be used as a label. Nor may any opcode or pseudo-op mnemonic, register name or condition code name. You should note the additional condition code names described later.

12.3.1.1. Temporary Labels

The assembler implements a system of temporary labels, useful for use within a localized section of code. These help eliminate the need to generate names for labels which are referenced only in the immediate vicinity of their definition, for example where a loop is implemented.

A temporary label takes the form of a digit string. A reference to such a label requires the same digit string, plus an appended b or f to signify a backward or forward reference respectively. Here is an example of the use of such labels.

```
entry_point: ;This is referenced from far away
    ld    b,10
1:  dec   c
    jr    nz,2f ;if zero, branch forward to 2:
    ld    c,8
    djnz  1b    ;decrement and branch back to 1:
    jr    1f    ;this does not branch to the
                ;same label as the djnz
2:  call  fred  ;get here from the jr nz,2f
1:  ret      ;get here from the jr 1f
```

The digit string may be any positive decimal number 0 to 65535. A temporary label value may be re-used any number of times. Where a reference to e.g. 1b is made, this will reference the closest label 1: found by looking backwards from the current point in the file. Similarly 23f will reference the first label 23: found by looking forwards from the current point in the file.

12.3.2. Constants

Constants may be entered in one of the radices 2, 8, 10 or 16. The default is 10. Constants in the other radices may be denoted by a trailing character drawn from the following set:

Character	Radix	Name
B	2	binary
O	8	octal
Q	8	octal
o	8	octal
q	8	octal
H	16	hexadecimal
h	16	hexadecimal

Hexadecimal constants may also be specified in C style, for example LD A,0x21. Note that a lower case b may not be used to indicate a binary number, since lb is a backward reference to a temporary label 1:.

12.3.2.1. Character Constants

A character constant is a single character enclosed in single quotes ('). Multi character constants may be used only as an operand to a DEFM pseudo-op.

12.3.2.2. Floating Constants

A floating constant in the usual notation (e.g. 1.234 or 1234e-3) may be used as the operand to a DEFF pseudo-op.

12.3.2.3. Opcode Constants

Any z80 opcode may be used as a constant in an expression. The value of the opcode in this context will be the byte that the opcode would have assembled to if used in the normal way. If the opcode is a 2-byte opcode (CB or ED prefix byte) only the second byte of the opcode will be used. This is particularly useful when setting up jump vectors. For example:

```
ld  a,jp      ;a jump instruction
ld  (0),a     ;0 is jump to warm boot
ld  hl,boot   ;done here
ld  (1),hl
```

12.3.3. Expressions

Expressions are constructed largely as described in the "Z80 Assembly Language Handbook".

12.3.3.1. Operators

The following operators may be used in expressions:

Operator	Meaning
&	Bitwise AND
*	Multiplication
+	Addition
-	Subtraction

.and.	Bitwise AND
.eq.	Equality test
.gt.	Signed greater than
.high.	Hi byte of operand
.low.	Low byte of operand
.lt.	Signed less than
.mod.	Modulus
.not.	Bitwise complement
.or.	Bitwise or
.shl.	Shift left
.shr.	Shift right
.ult.	Unsigned less than
.ugt.	Unsigned greater than
.xor.	Exclusive or
/	Divison
<	Signed less than
=	Equality
>	Signed greater than
^	Bitwise or

Operators starting with a dot "." should be delimited by spaces, thus label .and. 1 is valid but label.and.1 is not.

12.3.3.2. Relocatability

Zas produces object code which is relocatable; this means that it is not necessary to specify assembly time where the code is to be located in memory. It is possible to do so, by use of the ORG pseudo-op, however the preferred approach is to use program sections or psects. A psect is a named section of the program, in which code or data may be defined at assembly time. All parts of a psect will be loaded contiguously into memory, even if they were defined in separate files, or in the same file but separated by code for another psect. For example, the following code will load some executable instructions into the psect named text, and some data bytes into the data psect.

```
    psect text, global

alabel:
    ld    hl, astring
    call  putit
    ld    hl, anotherstring

    psect data, global
astring:
    defm  'A string of chars'
    defb  0
anotherstring:
    defm  'Another string'
    defb  0

    psect text

putit:
    ld    a, (hl)
    or    a
    ret   z
    call  outchar
    inc   hl
    jr    putit
```

Note that even though the two blocks of code in the text psect are separated by a block in the data psect, the two text psect blocks will be contiguous when loaded by the linker. The instruction "ld hl, anotherstring" will fall through to the label "putit:" during execution. The actual location in memory of the two psects will be determined by the linker. See the linker manual for information on how psect addresses are determined.

A label defined in a psect is said to be relocatable, that is, its actual memory address is not determined at assembly time. Note that this does not apply if the label is in the default (unnamed) psect, or in a psect declared absolute (see the PSECT pseudo-op description below). Any labels declared in an absolute psect will be absolute, that is their address will be determined by the assembler.

With the version of ZAS supplied with version 7 or later of HI-TECH C, relocatable expressions may be combined freely in expressions. Older versions of ZAS allowed only limited arithmetic on relocatable expressions.

12.3.4. Pseudo-ops

The pseudo-ops are based on those described in the "Z80 Assembly Language Handbook", with some additions.

12.3.4.1. DEFB, DB

This pseudo-op should be followed by a comma-separated list of expressions, which will be assembled into sequential

byte locations. Each expression must have a value between -128 and 255 inclusive. DB can be used as a synonym for DEFB. Example:

```
DEFB 10, 20, 'a', 0FFH
DB   'hello world',13,10,0
```

12.3.4.2. DEFF

This pseudo-op assembles floating point constants into 32 bit HI-TECH C format floating point constants. For example:

```
pi: DEFF 3.14159
```

12.3.4.3. DEFW

This operates in a similar fashion to DEFB, except that it assembles expressions into words, without the value restriction. Example:

```
DEFW -1, 3664H, 'A', 3777Q
```

12.3.4.4. DEFS

DeFS reserves memory locations without initializing them. Its operand is an absolute expression, representing the number of bytes to be reserved. This expression is added to the current location counter. Note however that locations reserved by DEFS may be initialized to zero by the linker if the reserved locations are in the middle of the program. Example:

```
DEFS 20h ;reserve 32 bytes of memory
```

12.3.4.5. EQU

EQU sets the value of a symbol on the left of EQU to the expression on the right. It is illegal to set the value of a symbol which is already defined. Example:

```
SIZE      equ    46
```

12.3.4.6. DEFL

This is identical to EQU except that it may redefine existing symbols. Example:

```
SIZE      defl  48
```

12.3.4.7. DEFM

Defm should be followed by a string of characters, enclosed in single quotes. The ASCII values of these characters are assembled into successive memory locations. Example:

```
DEFM  'A string of funny *@$ characters'
```

12.3.4.8. END

The end of an assembly is signified by the end of the source file, or the END pseudo-op. The END pseudo-op may optionally be followed by an expression which will define the start address of the program. This is not actually useful for CP/M. Only one start address may be defined per program, and the linker will complain if there are more. Example:

```
END    somelabel
```

12.3.4.9. COND, IF, ELSE, ENDC

Conditional assembly is introduced by the COND pseudo-op. The operand to COND must be an absolute expression. If its value is false (zero) the code following the COND up to the corresponding ENDC pseudo-op will not be assembled. COND/ENDC pairs may be nested. IF may be used as a synonym for COND. The ELSE pseudo operation may be included within a COND/ENDC block, for example:

```
IF     CPM
call   5
ELSE
call   os_func
ENDC
```

12.3.4.10. ELSE

See COND.

12.3.4.11. ENDC

See COND.

12.3.4.12. ENDM

See MACRO.

12.3.4.13. PSECT

This pseudo-op allows specification of relocatable program sections. Its arguments are a psect name, optionally followed by a list of psect flags. The psect name is a symbol constructed according to the same rules as for labels, however a psect may have the same name as a label without conflict. Psect names are recognized only after a PSECT pseudo-op. The psect flags are as follows:

ABS	Psect is absolute
GLOBAL	Psect is global
LOCAL	Psect is not global
OVRD	Psect is to be overlapped by linker
PURE	Psect is to be read-only

If a psect is global, the linker will merge it with any other global psects of the same name from other modules. Local psects will be treated as distinct from any other psect from another module. Psects are global by default.

By default the linker concatenates code within a psect from various modules. If a psect is specified as OVRD, the linker will overlap each module's contribution to that psect. This is particularly useful when linking modules which initialize e.g. interrupt vectors.

The PURE flag instructs the linker that the psect is to be made read-only at run time. The usefulness of this flag depends on the ability of the linker to enforce the requirement. CP/M fails miserably in this regard.

The ABS flag makes a psect absolute. The psect will be loaded at zero. This is useful for statically initializing interrupt vectors and jump tables. Examples:

```
PSECT    text, global, pure
PSECT    data, global
PSECT    vectors, ovrd
```

12.3.4.14. GLOBAL

Global should be followed by one more symbols (comma separated) which will be treated by the assembler as global symbols, either internal or external depending on whether they are defined within the current module or not. Example:

```
GLOBAL      labell, putchar, _printf
```

12.3.4.15. ORG

An ORG pseudo-op sets the current psect to the default (absolute) psect, and the location counter to its operand, which must be an absolute expression. Example:

```
ORG      100H
```

12.3.4.16. MACRO

This pseudo-op defines a macro. It should be either preceded or followed by the macro name, then optionally followed by a comma-separated list of formal parameters. The lines of code following the MACRO pseudo-op up to the next ENDM pseudo-op will be stored as the body of the macro. The macro name may subsequently be used in the opcode part of an assembler statement, followed by actual parameters. The text of the body of the macro will be substituted at that point, with any use of the formal parameters substituted with the corresponding actual parameter. For example:

```
print      MACRO string
    psect data
999:       db      string,'$'
    psect text
    ld      de,999b
    ld      c,9
    call    5
    ENDM
```

When used, this macro will expand to the 3 instructions in the body of the macro, with the actual parameters substituted for func and arg. Thus

```
print 'hello world'
```

expands to

```
    psect data
999:       db      'hello world','$'
    psect text
    ld      de,999b
    ld      c,9
    call    5
```

Macro arguments can be enclosed in angle brackets ('<')

and '>') to pass arbitrary text including delimiter characters like commas as a single argument. For example, suppose you wanted to use the print macro defined above to print a string which includes the carriage return and linefeed characters. The macro invocation:

```
print 'hello world',13,10
```

would fail because 13 and 10 are treated as extra arguments and ignored. In order to pass a string which includes commas as a single argument, you could write:

```
print <'hello world',13,10>
```

which would cause the text 'hello world',13,10 to be passed through as a single argument. This would expand to the following code:

```
psect data
999:    db    'hello world',13,10,'$'
psect text
ld     de,999b
ld     c,9
call   5
```

ZAS supports two forms of macro declaration for compatibility with older versions of ZAS and other Z80 assemblers. The macro name may be declared either in the label field before the MACRO pseudo-op, or in the operand field after the MACRO pseudo-op. Thus these two MACRO declarations are equivalent:

```
bdos    MACRO func,arg
ld      de,arg
ld      c,func
call    5
ENDM
```

and

```
MACRO bdos,func,arg
ld      de,arg
ld      c,func
call    5
ENDM
```

12.3.4.17. LOCAL

The LOCAL pseudo-op allows unique labels to be defined for each expansion of a macro. Any symbols listed after the LOCAL directive will have a unique assembler-generated symbol substituted for them when the macro is expanded. For

example:

```
copy      MACRO source,dest,count
LOCAL nocopy
push af
push bc
ld bc,source
ld a,b
or c
jr z,nocopy
push de
push hl
ld de,dest
ld hl,source
ldir
pop hl
pop de
nocopy:   pop bc
pop af
ENDM
```

when expanded will include a unique assembler generated label in place of nocopy. For example, copy (recptr),buf,(recsize) will expand to:

```
push af
push bc
ld bc,(recsize)
ld a,b
or c
jr z,??0001
push de
push hl
ld de,buf
ld hl,(recptr)
ldir
pop hl
pop de
??0001:  pop bc
pop af
```

if invoked a second time, the label nocopy would expand to ??0002.

12.3.4.18. REPT

The REPT pseudo-op defines a temporary macro which is then expanded a number of times, as determined by its argument. For example:

```
REPT 3
ld (hl),0
inc hl
ENDM
```

will expand to


```

ld    (hl),0
inc   hl
ld    (hl),0
inc   hl
ld    (hl),0
inc   hl

```

12.3.5. IRP and IRPC

The IRP and IRPC directives are similar to REPT, however instead of repeating the block a fixed number of times it is repeated once for each member of an argument list. In the case of IRP the list is a conventional macro argument list, in the case of IRPC it is successive characters from a string. For example:

```

IRP   string,<'hello world',13,10>,'arg2'
LOCAL str
psect data
str:  db    string,'$'
psect text
ld    c,9
ld    de,str
call  5
ENDM

```

would expand to

```

psect data
??0001: db    'hello world',13,10,'$'
psect text
ld    c,9
ld    de,??0001
call  5
psect data
??0002: db    'arg2','$'
psect text
ld    c,9
ld    de,??0002
call  5

```

Note the use of LOCAL labels and angle brackets in the same manner as with conventional macros.

IRPC is best demonstrated using the following example:

```

IRPC  char,ABC
ld    c,2
ld    e,'char'
call  5
ENDM

```

will expand to:

```
ld    c,2
ld    e,'A'
call  5
ld    c,2
ld    e,'B'
call  5
ld    c,2
ld    e,'C'
call  5
```

12.3.6. Extended Condition Codes

The assembler recognizes several additional condition codes. These are:

Code	Equivalent	Meaning
alt	m	Arithmetic less than
llt	c	Logical less than
age	p	Arithmetic greater or equal
lge	nc	Logical greater or equal
di		Use after ld a,i for test-
ei		ing state of interrupt
		enable flag - enabled or
_d_i_s_a_b_l_e_d_r_e_s_p_e_c_t_i_v_e_l_y_.		

12.4. Assembler Directives

An assembler directive is a line in the source file which produces no code, but rather which modifies the behaviour of the assembler. Each directive is recognized by the presence of an asterisk in the first column of the line, followed immediately by a word, only the first character of which is looked at. the line containing the directive itself is never listed. The directives are:

*Title

Use the text following the directive as a title for the listing.

*Heading

Use the text following the directive as a subtitle for the listing; also causes an *Eject.

*List

May be followed by ON or OFF to turn listing on or off respectively. Note that this directive may be used inside a macro or include file to control listing of that macro or include file. The previous listing state will be restored on exit from the macro or include file.

***Include**

The file named following the directive will be included in the assembly at that point.

***Eject**

A new page will be started in the listing at that point. A form feed character in the source will have the same effect.

Some examples of the use of these directives:

*Title Widget Control Program

*Heading Initialization Phase

*Include widget.i

12.5. Diagnostics

An error message will be written on the standard error stream for each error encountered in the assembly. This message identifies the file name and line number and describes the error. In addition the line in the listing where the error occurred will be flagged with a single character to indicate the error. The characters and the corresponding messages are:

A: Absolute expression required

B: Bad arg to *L
Bad arg to IM
Bad bit number
Bad character constant
Bad jump condition

D: Directive not recognized
Digit out of range

E: EOF inside conditional
Expression error

G: Garbage after operands
Garbage on end of line

I: Index offset too large

J: Jump target out of range

L: Lexical error

M: Multiply defined symbol

O: Operand error

P: Phase error
Psect may not be local and global

R: Relocation error

S: Size error
Syntax error

U: Undefined symbol
Undefined temporary label
Unterminated string

12.6. Z80/Z180/64180 Instruction Set

The remainder of this chapter is devoted to a complete instruction set listing for the Z80, Z180, 64180 and NSC800 processors. The Z180 and 64180 will execute all Z80 instructions, although the timing is different.

13. Linker Reference Manual

HI-TECH C incorporates a relocating assembler and linker to permit separate compilation of C source files. This means that a program may be divided into several source files, each of which may be kept to a manageable size for ease of editing and compilation, then each object file compiled separately and finally all the object files linked together into a single executable program.

The assembler is described in the machine-specific manual. This appendix describes the theory behind and the usage of the linker.

13.1. Relocation and Psects

The fundamental task of the linker is to combine several relocatable object files into one. The object files are said to be relocatable since the files have sufficient information in them so that any references to program or data addresses (e.g. the address of a function) within the file may be adjusted according to where the file is ultimately located in memory after the linkage process. Thus the file is said to be relocatable. Relocation may take two basic forms; relocation by name, i.e. relocation by the ultimate value of a global symbol, or relocation by psect, i.e. relocation by the base address of a particular section of code, for example the section of code containing the actual executable instructions.

13.1.1. Program Sections

Any object file may contain bytes to be stored in memory in one or more program sections, which will be referred to as psects. These psects represent logical groupings of certain types of code bytes in the program. The section of the program containing executable instructions is normally referred to as the text psect. Other sections are the initialized data psect, called simply the data psect, and the uninitialized data psect, called the bss psect.

In fact the linker will handle any number of psects, and in fact more may be used in special applications. However the C compiler uses only the three mentioned, and the names text, data and bss are simply chosen for identification; the linker assigns no special significance to the name of a psect.

The difference between the data and bss psects may be exemplified by considering two external variables; one is initialized to the value 1, and the other is not initialized. The first will be placed into the data psect, and the second in the bss psect. The bss psect is always cleared to zeros on startup of the program, thus the second variable will be initialized at run time to zero. The first will however occupy space in the program file, and will maintain its

initialized value of 1 at startup. It is quite possible to modify the value of a variable in the data psect during execution, however it is better practice not to do so, since this leads to more consistent use of variables, and allows for restartable and romable programs.

The text psect is the section into which all executable instructions are placed. On CP/M-80 the text psect will normally start at the base of the TPA, which is where execution commences. The data psect will normally follow the text psect, and the bss will be last. The bss does not occupy space in the program (.COM) file. This ordering of psects may be overridden by an option to the linker. This is especially useful when producing code for special hardware.

For MS-DOS and CP/M-86 the psects are ordered in the same way, but since the 8086 processor has segment registers providing relocation, both the text and data psects start at 0, even though they will be loaded one after the other in memory. This allows 64k code and 64k data and stack. Sufficient information is placed in the executable file (.EXE or .CMD) for the operating system to load the program in memory.

13.1.2. Local Psects and the Large Model

Since for practical purposes the psects are limited to 64K on the 8086, to allow more than 64K code the compiler makes use of local psects. A psect is considered local if the .psect directive has a LOCAL flag. Any number of local psects may be linked from different modules without being combined even if they have the same name. Note however that no local psect may have the same name as a global psect.

All references to a local psect within the same module (or within the same library) will be treated as references to the same psect. Between modules however two local psects of the same name are treated as distinct. In order to allow collective referencing of local psects via the -P option (described later) a local psect may have a class name associated with it. This is achieved with the CLASS flag on the .psect directive.

13.2. Global Symbols

The linker handles only symbols which have been declared as global to the assembler. From the C source level, this means all names which have storage class external and which are not declared as static. These symbols may be referred to by modules other than the one in which they are defined. It is the linker's job to match up the definition of a global symbol with the references to it.

13.3. Operation

A command to the linker takes the following form:

LINK options files ...

Options is zero or more linker options, each of which modifies the behaviour of the linker in some way. Files is one or more object files, and zero or more library names. The options recognized by the linker are as follows: they will be recognized in upper or lower case.

- R Leave the output relocatable.
- L Retain absolute relocation info. -LM will retain only segment relocation information.
- I Ignore undefined symbols.
- N Sort symbols by address.
- Caddr
Produce a binary output file offset by addr.
- S Strip symbol information from the output file.
- X Suppress local symbols in the output file.
- Z Suppress trivial (compiler-generated) symbols in the output file.
- Oname
Call the output file name.
- Pspec
Spec is a psect location specification.
- Mname
Write a link map to the file name.
- Usymbol
Make symbol initially undefined.
- Dfile
Write a symbol file.
- Wwidth
Specify map width.

Taking each of these in turn:

The -R option will instruct the linker to leave the output file (as named by a -O option, or l.obj by default) relocatable. This is normally because there are further files to be linked in, and the output of this link will be used as input to the linker subsequently. Without this option, the linker will make the output file absolute, that

is with all relocatable addresses made into absolute references. This option may not be used with the -L or -C options.

The -L option will cause the linker to output null relocation information even though the file will be absolute. This information allows self-relocating programs to know what addresses must be relocated at run time. This option is not usable with the -C option. In order to create an executable file (i.e. a .COM file) the program objtohex must be used. If a -LM option is used, only segment relocation information will be retained. This is used in conjunction with the large memory model. Objtohex will use the relocation information (when invoked with a -L flag) to insert segment relocation addresses into the executable file.

The -I option is used when it is desired to link code which contains symbols which are not defined in any module. This is normally only used during top-down program development, when routines are referenced in code written before the routines themselves have been coded.

When obtaining a link map via the -M option, the symbol table is by default sorted in order of symbol name. To sort in order of address, the -N option may be used.

The output of the linker is by default an object file. To create an executable program, this must be converted into an executable image. For CP/M this is a .COM file, which is simply an image of the executable program as it should appear in memory, starting at location 100H. The linker will produce such a file with the -C100H option. File formats for other applications requiring an image binary file may also be produced with the -C option. The address following the -C may be given in decimal (default), octal (by using o or O suffix) or hexadecimal (by using an h or H suffix).

Note that because of the complexity of the executable file formats for MS-DOS and CP/M-86, LINK will not produce these (.EXE and .CMD resp.) formats directly. The compiler automatically runs OBJTOHEX with appropriate options to generate the correct file format.

The -S, -X and -Z options, which are meaningless when the -C option is used, will strip respectively all symbols, all local symbols or all trivial local symbols from the output file. Trivial symbols are symbols produced by the compiler, and have the form of one of a set of alphabetic characters followed by a digit string.

The default output file name is l.obj, or l.bin when the -C option is used. This may be overridden by the -Oname option. The output file will be called name in this instance. Note that no suffix is appended to the name; the file will be called exactly the argument to the option.

For certain specialized applications, e.g. producing

code for an embedded microprocessor, it is necessary to specify to the linker at what address the various psects should be located. This is accomplished with the -P option. It is followed by a specification consisting of a comma-separated list of psect names, each with an optional address specification. In the absence of an address specification for a psect listed, it will be concatenated with the previous psect. For example

```
-Ptext=0c000h,data,bss=8000h
```

This will cause the text psect to be located at 0C000H, the data psect to start at the end of the text psect, and the bss psect to start at 8000H. This may be for a processor with ROM at 0C000H and RAM at 8000H.

Where the link address, that is the address at which the code will be addressed at execution time, and the load address, that is the address offset within the output file, are different (e.g for the 8086) it is possible to specify the load address separately from the link address. For example:

```
-Ptext=100h/0,data=0C000h/
```

This specification will cause the text segment to be linked for execution at 100h, but loaded in the output file at 0, while the data segment will be linked for 0C000h, but loaded contiguously with the text psect in the file. Note that if the slash ('/') is omitted, the load address is the same as the link address, while if the slash is supplied, but not followed by an address, the psect will be loaded after the previous psect.

In order to specify link and load addresses for local psects, the group name to which the psects belong may be used in place of a global psect name. The local psects will then have a link address as specified in the -P option, and load addresses incrementing upwards from the specified load address.

The -Mname option requests a link map, containing symbol table and module load address information to be written onto the file name. If name is omitted, the map will be written to standard output. -W may be used to specify the desired width of the map.

The -U option allows the specification to the linker of a symbol which is to be initially entered into the symbol table as undefined. This is useful when loading entirely from libraries. More than one -U flag may be used.

If it is desired to use the debugger on the program being linked, it is useful to produce a symbol file. The -Dfile option will write such a symbol file onto the named file, or l.sym if no file is given. The symbol file consists

of a list of addresses and symbols, one per line.

13.4. Examples

Here are some examples of using the linker. Note however that in the normal case it is not necessary to invoke the linker explicitly, since it is invoked automatically by the C command.

```
LINK -MMAP -C100H START.OBJ MAIN.OBJ A:LIBC.LIB
```

This command links the files start.obj and main.obj with the library a:libc.lib. Only those modules that are required from the library will be in fact linked in. The output is to be in .COM format, placed in the default file l.bin. A map is to be written to the file of the name map. Note that the file start.obj should contain startup code, and in fact the lowest address code in that file will be executed when the program is run, since it will be at 100H.

```
LINK -X -R -OX.OBJ FILE1.OBJ FILE2.OBJ A:LIBC.LIB
```

The files file1.obj and file2.obj will be linked with any necessary routines from a:libc.lib and left in the file x.obj. This file will remain relocatable. Undefined symbols will not cause an error. The file x.obj will probably later be the object of another link invocation. All local symbols will be stripped from the output file, thus saving space.

13.5. Invoking the Linker

The linker is called LINK, and normally resides on the A: drive, under CP/M, or in the directory A:\HITECH\ under MS-DOS. It may be invoked with no arguments, in which case it will prompt for input from standard input. If the standard input is a file, no prompts will be printed. The input supplied in this manner may contain lower case, whereas CP/M converts the entire command line to upper case by default. This is useful with the -U and -P options. This manner of invocation is generally useful if the number of arguments to LINK is large. Even if the list of files is too long to fit on one line, continuation lines may be included by leaving a backslash ('\') at the end of the preceding line. In this fashion, LINK commands of almost unlimited length may be issued.

14. Librarian

The librarian program, LIBR, has the function of combining several object files into a single file known as a library. The purposes of combining several such object modules are several.

- a. fewer files to link
- b. faster access
- c. uses less disk space

In order to make the library concept useful, it is necessary for the linker to treat modules in a library differently from object files. If an object file is specified to the linker, it will be linked into the final linked module. A module in a library, however, will only be linked in if it defines one or more symbols previously known, but not defined, to the linker. Thus modules in a library will be linked only if required. Since the choice of modules to link is made on the first pass of the linker, and the library is searched in a linear fashion, it is possible to order the modules in a library to produce special effects when linking. More will be said about this later.

14.1. The Library Format

The modules in a library are basically just concatenated, but at the beginning of a library is maintained a directory of the modules and symbols in the library. Since this directory is smaller than the sum of the modules, the linker is speeded up when searching a library since it need read only the directory and not all the modules on the first pass. On the second pass it need read only those modules which are required, seeking over the others. This all minimizes disk i/o when linking.

It should be noted that the library format is geared exclusively toward object modules, and is not a general purpose archiving mechanism as is used by some other compiler systems. This has the advantage that the format may be optimized toward speeding up the linkage process.

14.2. Using

The librarian program is called LIBR, and the format of commands to it is as follows:

```
LIBR k file.lib file.obj ...
```

Interpreting this, LIBR is the name of the program, k is a key letter denoting the function requested of the librarian (replacing, extracting or deleting modules, listing modules or symbols), file.lib is the name of the library file to be operated on, and file.obj is zero or more object

file names.

The key letters are:

r	replace modules
d	delete modules
x	extract modules
m	list module names
s	list modules with symbols

When replacing or extracting modules, the file.obj arguments are the names of the modules to be replaced or extracted. If no such arguments are supplied, all the modules in the library will be replaced or extracted respectively. Adding a file to a library is performed by requesting the librarian to replace it in the library. Since it is not present, the module will be appended to the library. If the r key is used and the library does not exist, it will be created.

Under the d keyletter, the named object files will be deleted from the library. In this instance, it is an error not to give any object file names.

The m and s keyletters will list the named modules and, in the case of the s keyletter, the symbols defined or referenced within (global symbols only are handled by the librarian). As with the r and x keyletters, an empty list of modules means all the modules in the library.

14.3. Examples

Here are some examples of usage of the librarian.

LIBR m file.lib

List all modules in the library file.lib.

LIBR s file.lib a.obj b.obj c.obj

List the global symbols in the modules a.obj, b.obj and c.obj

LIBR r file.lib 1.obj 2.obj

Replace the module 1.obj in the file file.lib with the contents of the object file 1.obj, and repeat for 2.obj. If the object module is not already present in the library, append it to the end.

LIBR x file.lib

Extract, without deletion, all the modules in file.lib and write them as object files on disk.

LIBR d file.lib a.obj b.obj 2.obj

Delete the object modules a.obj, b.obj and 2.obj from the library file.lib.

14.4. Supplying Arguments

Since it is often necessary to supply many object file arguments to LIBR, and command lines are restricted to 127 characters by CP/M and MS-DOS, LIBR will accept commands from standard input if no command line arguments are given. If the standard input is attached to the console, LIBR will prompt. Multiple line input may be given by using a backslash as a continuation character on the end of a line. If standard input is redirected from a file, LIBR will take input from the file, without prompting. For example:

```
LIBR
libr> r file.lib 1.obj 2.obj 3.obj \
libr> 4.obj 5.obj 6.obj
```

will perform much the same as if the .obj files had been typed on the command line. The libr> prompts were printed by LIBR itself, the remainder of the text was typed as input.

```
LIBR <lib.cmd
```

Libr will read input from lib.cmd, and execute the command found therein. This allows a virtually unlimited length command to be given to LIBR.

14.5. Listing Format

A request to LIBR to list module names will simply produce a list of names, one per line, on standard output. The s keyletter will produce the same, with a list of symbols after each module name. Each symbol will be preceded by the letter D or U, representing a definition or reference to the symbol respectively. The -W option may be used to determine the width of the paper for this operation. For example LIBR -w80 s file.lib will list all modules in file.lib with their global symbols, with the output formatted for an 80 column printer or display.

14.6. Ordering of Libraries

The librarian creates libraries with the modules in the order in which they were given on the command line. When updating a library the order of the modules is preserved. Any new modules added to a library after it has been created will be appended to the end.

The ordering of the modules in a library is significant to the linker. If a library contains a module which references a symbol defined in another module in the same library, the module defining the symbol should come after the module referencing the symbol.

14.7. Error Messages

Libr issues various error messages, most of which represent a fatal error, while some represent a harmless occurrence which will nonetheless be reported unless the `-w` option was used. In this case all warning messages will be suppressed.

15. Objtohex

The HI-TECH linker is capable of producing simple binary files, or object files as output. Any other format required must be produced by running the utility program OBJTOHEX. This allows conversion of object files as produced by the linker into a variety of different formats, including various hex formats. The program is invoked thus:

```
OBJTOHEX options inputfile outputfile
```

All of the arguments are optional. If outputfile is omitted it defaults to l.hex or l.bin depending on whether the -b option is used. The inputfile defaults to l.obj.

The options are:

-Baddr

Produce a binary image output. This is similar to the -C option of the linker. If addr is supplied, the start of the image file will be offset by addr. If addr is omitted, the first byte in the file will be the lowest byte initialized. Addr may be given in decimal, octal or hexadecimal. The default radix is decimal, and suffix letters of o or O indicate octal, and h or H indicate hex. Thus -B100H will produce a file in .COM format.

-I Include symbol records in the Intel format hex output. Each symbol record has a form similar to an object record, but with a different record type. The data bytes in the record are the symbol name, and the address is the value of the symbol. This is useful for downloading to ROM debuggers.

-C Read a checksum specification from the standard input. The checksum specification is described below. Typically the specification will be in a file.

-Estack

This option produces an MS-DOS .EXE format file. The optional stack argument will determine the maximum stack size the program will be allocated on execution. By default the program will be allocated the maximum stack available, up to the limit of 64K data. If a stack argument is supplied, the stack size will not exceed the argument. This is useful to limit the amount of memory a program will use. The stack argument takes the same form as the argument to -B above.

-8stack

This option will produce a CP/M-86 .CMD file. The stack argument is the same as for the -E option.

-Astack

This is used when producing a.out format files for unix systems (specifically Venix-86). If the stack argument is zero, the size of the data segment will be 64k, otherwise the stack will be placed below the data segment, and its size set to stack. This must be co-ordinated with appropriate arguments to the -p option of the linker.

-M This flag will instruct objtohex to produce Motorola 'S' format hex output.

-L This option is used when producing large model programs; the linker will have been used with the -LM option to retain segment relocation information in the object file. Use of the -L option to objtohex will cause it to convert that segment relocation information into appropriate data in the executable file for use when the program is loaded. Either the operating system or the run-time startup code will use the relocation data to adjust segment references based on where in memory the program is actually loaded. If the -L option is followed by a symbol name, then the relocation information will be stored at the address represented by that symbol in the output file, e.g. -L__Bss will cause it to be stored at the base of the bss psect (__Bss is defined by the linker to be the load address of the bss psect). If the special symbol Dos_hdr is used then the relocation information will be stored in the .EXE file header. This is only valid in conjunction with the -E option.

-S The -S option instructs objtohex to write a symbol file. The symbol file name is given after the -S, e.g. -Sxx.sym.

Unless another format is specifically requested, objtohex will produce a file in Intel hex format. This is suitable for down-line loading, PROM programming etc. The HP format is useful for transferring code to an HP64000 for emulation or PROM programming.

The checksum specification allows automated checksum calculation. The checksum specification takes the form of several lines, each line describing one checksum. The syntax of a checksum line is:

addr1-addr2 where1-where2 +offset

All of addr1, addr2, where1, where2 and offset are hex numbers, without the usual H suffix. Such a specification says that the bytes at addr1 through to addr2 inclusive should be summed and the sum placed in the locations where1 through where2 inclusive. For an 8 bit checksum these two addresses should be the same. For a checksum stored low byte first, where1 should be less than where2, and vice versa. The +offset is optional, but if supplied, the value offset

will be used to initialize the checksum. Otherwise it is initialized to zero. For example:

```
0005-1FFF 3-4 +1FFF
```

This will sum the bytes in 5 through 1FFFH inclusive, then add 1FFFH to the sum. The 16 bit checksum will be placed in locations 3 and 4, low byte in 3. The checksum is initialized with 1FFFH to provide protection against an all zero rom, or a rom misplaced in memory. A run time check of this checksum would add the last address of the rom being checksummed into the checksum. For the rom in question, this should be 1FFFH. The initialization value may, however, be used in any desired fashion.

Use this page for notes

16. Cref

The cross reference list utility CREF is used to format raw cross-reference information produced by the compiler or the assembler into a sorted listing. A raw cross-reference file is produced with the -CR option to the compiler. The assembler will generate a raw cross-reference file with a -C option (Z80 or 8086 assemblers) or by using an OPT CRE directive (6800 series assemblers) or a REF control line (8096 assembler).. The general form of the CREF command is:

CREF options files

where options is zero or more options as described below and files is one or more raw cross-reference files. CREF takes the following options:

-Ooutfile

Allows specification of the output file name. By default the listing will be written to the standard output and may be redirected in the usual manner. Alternatively using the -O option an output file name may be specified, e.g. -Oxxx.lst.

-Pwidth

This option allows the specification of the width to which the listing is to be formatted, e.g. -P132 will format the listing for a 132 column printer. The default is 80 columns.

-Llength

Specify the length of the paper on which the listing is to be produced, e.g. if the listing is to be printed on 55 line paper you would use a -L55 option. The default is 66 lines.

-Xprefix

The -X option allows the exclusion of symbols from the listing, based on a prefix given as argument to -X. For example if it was desired to exclude all symbols starting with the character sequence xyz then the option -Xxyz would be used. If a digit appears in the character sequence then this will match any digit in the symbol, e.g. -XX0 would exclude any symbols starting with the letter X followed by a digit.

-F -F will exclude from the listing any references from files with a full path name. A full path name means either: a file name starting with a slash ('/') or backslash ('\') or a file name starting with a CP/M user number/drive letter prefix, e.g. 0:A:. This is intended to force omission from the listing of any symbol references derived from standard header files, e.g. using -F would omit any references from the header file STDIO.H.

-Hstring

The -H option takes a string as an argument which will be used as a header in the listing. The default heading is the name of the first raw cross-ref information file specified.

-Sstoplist

The -S option should have as its argument the name of a file containing a list of symbols not to be listed in the cross-reference. Multiple stoplists may be supplied with multiple -S options.

Cref will accept wild card filenames and I/O redirection. Long command lines may be supplied by invoking CREF with no arguments and typing the command line in response to the cref> prompt. A backslash at the end of the line will be interpreted to mean that more command lines follow.

APPENDIX 1

Error Messages

Error Messages produced by the compiler are listed below. Each message is followed by the name of the program which produces it, and some further description of what causes the message or what to do about it.

'.' expected after '...' P1
The ellipsis symbol must have three dots

actuals too long CPP
Reduce length of macro arguments

argument list conflicts with prototype P1
The argument list in a function definition must agree with a prototype if one exists

argument redeclared P1
This argument has been declared twice

arithmetic overflow in constant expression CGEN
Evaluation of this constant expression produced an arithmetic overflow. This may or may not represent a true error.

array index out of bounds P1
An array index expression evaluates to a constant which is less than zero or greater than or equal to the dimension of the array

Assertion CGEN
Internal error - contact HI-TECH

attempt to modify const object P1
An attempt has been made to assign to or otherwise modify an object designated as 'const'

bad bitfield type P1
Bitfields must be of type 'int'

Bad conval CGEN
Internal error - contact HI-TECH

Bad dimensions CGEN
An array has bad dimensions - probably zero

Bad element count expr CGEN
Internal error - contact HI-TECH

bad formal CPP
Check macro definition syntax

bad include syntax CPP
 Use only "" and <> for include files

Bad int. code CGEN
 The intermediate code file has been corrupted - can be
 caused by running out of disk space

Bad -M option CGEN
 A -M option passed to the code generator is unknown

Bad mod '+' for how = c CGEN
 Internal error - contact HI-TECH

bad object code format LINK
 This file is either corrupted or not a valid object
 file

Bad op d to swaplog CGEN
 Internal error - contact HI-TECH

Bad op n to revlog CGEN
 Internal error - contact HI-TECH

bad origin format in spec LINK
 An address in a -p option is invalid

bad '-p' format LINK
 The -p option provided is invalid

Bad pragma c CGEN
 The code generator has been passed a pragma it does not
 know about

Bad putwsize CGEN
 Internal error - contact HI-TECH

bad storage class P1, CGEN
 The specified storage class is illegal

Bad U usage CGEN
 Internal error - contact HI-TECH

Bit field too large (n bits) CGEN
 A bit field may not be larger than an int

Cannot get memory LINK
 The linker has run out of dynamic memory

Can't be both far and near P1
 The 'far' and 'near' keywords cannot appear in the same
 type specifier

can't be long P1
 Chars and shorts cannot be long

can't be register P1
 An extern or static variable may not be register

can't be short P1
Float and char cannot be short

can't be unsigned P1
Float cannot be unsigned

can't call an interrupt function P1
A function qualified 'interrupt' can only be called by hardware, not by an ordinary function call

Can't create filename CGEN
The file specified could not be created

Can't create xref file P1
The cross reference file specified could not be created

Can't create CPP
Output file could not be created

Can't create LINK
The linker cannot create a file

Can't find include file CPP
Check and correct the include file name - spaces are not allowed in file names

Can't find register for bits CGEN
Internal error - contact HI-TECH

Can't generate code for this expression CGEN
The code generator is unable to generate code for this expression - simplifying the expression (e.g. computing values into temporary variables) will usually correct it, otherwise contact HI-TECH

can't have array of functions P1
You cannot have an array of functions - you can have an array of pointers to functions

Can't have 'port' variable CGEN
You cannot declare a variable to be qualified 'port' - you can only use port to qualify pointers or typecast constant values

can't have storage class P1
A storage class may not appear in a prototype argument

can't initialise auto aggregates P1
You cannot initialise a structure or array inside a function unless it is static

can't initialize arg P1
An argument cannot have an initializer

can't mix proto and non-proto args P1
You cannot mix prototype and non-prototype arguments even in a function definition

Can't open filename CGEN
The file specified could not be opened for reading

Can't open LINK
The linker cannot open a file

Can't seek LINK
The linker could not seek in file

can't take address of register variable P1
You can't take the address of a variable in a register

can't take sizeof func CGEN
You can't take the size of a function. You can take the size of a function call

can't take this address P1
The expression does not have an address

'case' not in switch P1
A 'case' label is permitted only inside a switch

char const too long P1
A character constant may have only one character in it

close error (disk space?) P1
Probably out of disk space

common symbol psect conflict LINK
A common symbol is defined to be in more than one psect

constant conditional branch CGEN
You have a program structure testing a constant expression, e.g. while(1). You should substitute for this the more efficient for(;;)

constant expression required P1
A constant expression is required in e.g. an array dimension

constant operand to || or && CGEN
A logical operator has a constant operand which has been optimized out

declarator too complex P1
This declaration is too complex for the compiler to handle

default case redefined P1
Only one default case is permitted in a switch

'default' not in switch P1
A 'default' label is permitted only inside a switch

digit out of range P1
An octal constant may not contain 7 or 8, and a decimal constant may not contain A-F

dimension required P1
A dimension is required for all except the most significant in an array declaration

Division by zero CGEN
Attempt to divide by zero in this expression

Duplicate case label n CGEN
There are two case labels in this switch that have the same value

Duplicate -d flag LINK
Only one -d flag is allowed to the linker

duplicate label P1
This label is defined twice

Duplicate -m flag LINK
Only one -m flag is allowed to the linker

duplicate qualifier P1
The same qualifier appears more than once in this type specifier

entry point multiply defined LINK
A program can only have one entry point (start address)

EOF in #asm P1
End of file was encountered after #asm and before a #endasm was seen

Error closing output file CGEN, CPP
Probably means you have run out of disk space

excessive -I file ignored CPP
Use fewer -I options

expand - bad how CGEN
Internal error - contact HI-TECH

expand - bad which CGEN
Internal error - contact HI-TECH

exponent expected P1
An exponent is expected after the 'e' or 'E' in a floating point constant. The exponent must contain only +, - and digits 0-9

Expression error CGEN
Internal error - contact HI-TECH

expression generates no code CGEN
This expression has no side effects and thus generates no code. It has been optimized out

expression syntax P1
The expression is badly formed

expression too complex P1
The expression has too many nested parantheses or other nested constructs

Fixup overflow referencing LINK
The linker has relocated a reference to a psect or symbol and the relocated address is too big to fit into the space, e.g. a relocated one byte address exceeds 256 or a relocated 16 bit address exceeds 65536

float param coerced to double P1
This float parameter has been converted to double - a prototype will override this coercion

function() declared implicit int P1
This function has been called without an explicit declaration. It is wise to explicitly declare all functions, preferably with a prototype. This will avoid many potential errors where your program comprises more than one source file

function does not take arguments P1
The prototype for this function indicates it takes no arguments

function or function pointer required P1
A function identifier or pointer to function is required for a function call.

functions can't return arrays P1
A function cannot return an array - it can return a pointer

functions can't return functions P1
A function cannot return a function - it can return a pointer to function

hex digit expected P1
A hex digit is expected after '0x'

identifier is a structure tag P1
A structure tag has been used in a context where another kind of tag is expected, e.g. saying struct fred where fred has previously been declared as union fred.

identifier is a union tag P1
Similar to the above error

identifier is an enum tag P1
Similar to the above error

identifier: large offset CGEN
Z80 only: This identifier has a large offset from the stack frame and thus access to it is inefficient. In a function any arrays should be declared after any simple variables

identifier redeclared P1
 The identifier has been redeclared with different attributes

identifier redefined P1
 An identifier has been defined twice

If-less else CPP
 Check #if usage

If-less endif CPP
 Check #if usage

illegal '#' directive P1
 A # directive passed through to the first pass is unknown. If this occurs with a #include it may be caused by a previous include file not having a <CR><LF> or newline on the last line.

Illegal character in preprocessor if CPP
 Check for strange character

illegal character P1
 A character unknown to the compiler has been encountered. The value given is the octal value of the character

illegal conversion between pointer types P1
 The expression causes one pointer type to be converted to another incompatible type

illegal conversion of integer to pointer P1
 An integer is used where a pointer is expected

illegal conversion of pointer to integer P1
 A pointer is used where an integer is expected

illegal conversion P1
 The type conversion here is illegal

Illegal flag LINK
 This option is illegal

illegal function qualifier(s) P1
 A function cannot have 'const' qualification

illegal initialisation P1
 The initialisation of this variable is illegal

Illegal number CPP
 Check number syntax

illegal type for array dimension P1
 An array dimension must be an integral quantity

illegal type for index expression P1
 An array index must be a simple integral expression

illegal type for switch expression P1
The expression in a 'switch' must be integral

illegal use of void expression P1
Void expressions may not be used in any way

implicit conversion of float to integer P1
A floating point value has been converted to integer -
truncation may occur

implicit return at end of non-void function P1
A function with a non-void type has returned without a
return statement

implicit signed to unsigned conversion P1
Unwanted sign extension may occur here. Add an explicit
typecast to force exactly the conversion you want

inappropriate break/continue P1

inappropriate 'else' P1
An 'else' has appeared without a matching 'if'

inconsistent storage class P1
Only one storage class may be specified in a declara-
tion

inconsistent type P1
Only one basic type may be specified in a declaration

initialisation illegal in arg list P1
You cannot initialise a function parameter

initialisation syntax P1
The syntax of this initialisation is illegal

initializer in 'extern' declaration P1
A declaration with the 'extern' keyword has an initial-
izer; this is not permitted as the extern declaration
reserves no storage

integer constant expected P1
An integer constant was expected here

integer expression required P1
An integral expression is required here

integral type required P1
An integral type is required here

large offset CGEN
Z80 only: This identifier has a large offset from the
stack frame and thus access to it is inefficient. In a
function any arrays should be declared after any simple
variables

Line too long P1
The source line is too long, or does not have a
<CR><LF> or newline at the end

local psect conflicts with global psect of same name
LINK
A local psect cannot have the same name as a global
psect

logical type required P1
A logical type (i.e. an integral type) is required as
the subject of a conditional expression

lvalue required P1
An lvalue, i.e. something which can be assigned to, is
required after an '&' or on the left hand of an assign-
ment

macro recursion CPP
A preprocessor macro has attempted to expand itself.
This would create infinite recursion

member is not a member of the struct/union P1
This member is not in the structure or union with which
it is used

members cannot be functions P1
A member cannot be a function - it can be a pointer to
function

Missing arg to -u LINK
-u requires an argument

Missing arg to -w LINK
-w requires an argument

missing) CPP
Put correct) in expression

Missing number after % in -p option LINK
After % in a -p option there must be a number

Missing number after pragma 'pack' P1
The correct syntax is #pragma pack(n) where n is 1, 2
or 4.

module has code below file base LINK
A -C option was specified but the program has code
below the address specified as the base of the binary
file

multiply defined symbol LINK
A symbol is defined more than once

name is a union, struct or enum P1
A union, struct or enum tag has been re-used in a dif-
ferent context

No case labels CGEN
This switch has no case labels

no identifier in declaration P1
This declaration should have an identifier in it

No room CGEN
The code generator has run out of dynamic memory. You will need to reduce the number of symbols and/or the complexity of expressions

No source file CPP
Source file could not be found - check spelling, directory paths etc.

no space CPP
Reduce number/size of macro definitions

no start record: entry point defaults to zero LINK
No start address has been specified for the program; the linker has set the start address to 0

Non-constant case label CGEN
This case label does not evaluate to an integral constant

non-void function returns no value P1
A function which should return a value has a 'return' statement with no value

not a variable identifier P1
The identifier is not a variable - it may be e.g. a label or structure tag

not an argument P1
This identifier is not in the argument list for this function

only functions may be qualified interrupt P1
The type qualifier 'interrupt' may be applied only to functions, not variables.

only functions may be void P1
Only functions, not variables, may be declared void

only lvalues may be assigned to or modified P1
You have attempted to modify an expression which does not identify a storage location

only register storage class allowed P1
A parameter may only be auto or register

operands of operator not same pointer type P1
The operands to the named operator in the expression are both pointers but are not the same pointer type

operands of operator not same type P1
The operands to the named operator in the expression
are incompatible types

pointer required P1
A pointer is required after a '*' (indirection) opera-
tor

popreg - bad reg CGEN
Internal error - contact HI-TECH

portion of expression has no effect CGEN
A portion of this expression has no effect on its value
and no side effects

probable missing '}' in previous block P1
A declaration has been encountered where an expression
was expected. The likely cause of this is that you have
omitted a closing '}' in the function above this point.

psect cannot be in classes a and b LINK
A psect can only be in one class

psect exceeds max size LINK
This psect is larger than a specified maximum size

psect is absolute LINK
This psect is absolute and cannot have a link address
specified in a -p option

Psect not loaded on 0xhexnum boundary LINK
This psect must be loaded on a specific boundary

Psect not relocated on 0xhexnum boundary LINK
This psect must be linked on a specific boundary

psect origin multiply defined LINK
This psect has its link address defined more than once

pushreg - bad reg CGEN
Internal error - contact HI-TECH

redundant & applied to array P1
An array type has an '&' operator applied to it. It has
been ignored since use of an array implicitly gives its
address

regused - bad arg to G CGEN
Internal error - contact HI-TECH

signatures do not matchLINK
An extern function has been declared with an incorrect
prototype. For example if an argument is declared as a
long in an extern declaration, but is really an int, a
signature mismatch will occur.

signed bitfields not supported P1
Only unsigned bitfields are supported

simple type required P1
An array or structure type cannot be used here

Sizeof yields 0 CGEN
The size of an object has evaluated to zero in a context where this is illegal, e.g. incrementing a pointer to a zero length object.

storage class illegal P1
A storage class may not be specified here

struct/union member expected P1
A structure or union member is required after a '.' or '->'

struct/union redefined P1
This structure or union has been defined twice

struct/union required P1
A structure or union identifier is required before a '.'

Switch on long! CGEN
Switching on a long expression is not supported

symbol cannot be global LINK
Stack, filename or line number symbols cannot be global

Syntax error in checksum list LINK
The checksum list provided is invalid

token too long CPP
Shorten token (e.g. identifier)

too few arguments P1
The prototype for this function lists more arguments than have been supplied

too many arguments P1
More arguments have been supplied than listed in the prototype for this function

Too many cases in switch CGEN, P1
There are too many cases in this switch

too many -D options CPP
Use fewer -D options

too many defines CPP
Reduce number of macro definitions

Too many errors CGEN
CGEN has given up because there were too many errors.

too many formals CPP
 Reduce number of parameters to this macro definition

Too many initializers CGEN
 There are too many initializers for this object

Too many psects LINK
 There are too many psects for the symbol table

Too many symbols LINK
 There are too many symbols for the linker symbol table

too many -U options CPP
 Use fewer -U options

too much defining CPP
 Reduce number/size of macros

too much indirection P1
 Too many '*'s in this declaration

too much pushback CPP
 Simplify macro usage

type conflict P1
 There is a conflict of types in this expression, e.g.
 attempting to assign a structure to a simple type

type specifier reqd. for proto arg P1
 A prototype argument must have a basic type

undefined control CPP
 Check use of #

undefined enum tag P1
 This enumerated type tag has not been defined

undefined identifier P1
 This identifier has not been defined before use

undefined struct/union P1
 The structure or union used has not been defined

undefined symbol LINK
 A list of undefined symbols follows. If some of the
 symbols should be in a library which was linked, it may
 be caused by a library ordering problem. In this case
 rebuild the library with the correct ordering or
 specify the library more than once in the link command

unexpected EOF P1
 End of file was encountered in the middle of a C con-
 struct. This is commonly caused by omission of a clos-
 ing '}' earlier in the program.

Unknown predicate CGEN
 Internal error - contact HI-TECH

unknown psect LINK

The psect specified in a -p option is not present in the program. Check the spelling and check the case - upper case does not match lower case

unreachable code P1

This section of code can never be executed as there is no possible path to reach it

Unreasonable include nesting CPP

Reduce number of include files

Unreasonable matching depth CGEN

Internal error - contact HI-TECH

unterminated macro call CPP

Probably missing)

void function cannot return value P1

A function declared void cannot return a value

Write error (out of disk space?) LINK

Probably means the disk is full

APPENDIX 2

Standard Library Functions

The functions accessible to user programs in the standard library `libc.lib` are listed below, by category with a short comment, then alphabetically with a longer description. In the detailed description of each function, the SYNOPSIS section describes the function in roughly the manner in which the function would be declared in the source file defining it. Where an include file is shown, this implies that that include file must be included in any source file using that function.

Where an include file is not provided, it will normally be necessary for an extern declaration of the function to be included in any source module using it, to ensure that the type of the function is correct. For example, if the function `lseek()` was to be used, a declaration of the form

```
extern long lseek();
```

should be in either the source file itself or an include file included in the source file. This ensures that the compiler knows that `lseek()` returns a long value and not the default `int`.

Where reference is made to `STDIO`, this means the group of functions under the heading STANDARD I/O below. These all have one thing in common; they operate on pointers to a defined data type called `FILE`. Such a pointer is often referred to as a stream pointer. The concept of a stream is central to these routines. Essentially a stream is a source or sink of data bytes. To the operating system and library routines this stream is featureless, i.e. no record structure is implied or assumed. Some routines do however recognize end of line characters.

STANDARD I/O

<code>fopen(name, mode)</code>	Open file for I/O
<code>freopen(name, mode, stream)</code>	Re-open existing stream
<code>fdopen(fd, mode)</code>	Associate a stream with a file descriptor
<code>fclose(stream)</code>	Close open file
<code>fflush(stream)</code>	Flush buffered data
<code>getc(stream)</code>	Read byte from stream
<code>fgetc(stream)</code>	Same as <code>getc</code>
<code>ungetc(c, stream)</code>	Push char back onto stream
<code>putc(c, stream)</code>	Write byte to stream
<code>fputc(c, stream)</code>	Same as <code>putc()</code>
<code>getchar()</code>	Read byte from standard input
<code>putchar(c)</code>	Write byte to standard output
<code>getw(stream)</code>	Read word from stream

putw(w, stream)	Write word to stream
gets(s)	Read line from standard input
fgets(s, n, stream)	Read string from stream
puts(s)	Write string to standard output
fputs(s, stream)	Write string to stream
fread(buf, size, cnt, stream)	Binary read from stream
fwrite(buf, size, cnt, stream)	Binary write to stream
fseek(stream, offs, wh)	Random access positioning
ftell(stream)	Current file read/write position
rewind(stream)	Reposition file pointer to start
setvbuf(stream, buf, mode, size)	Enable/disable buffering of stream
fprintf(stream, fmt, args)	Formatted output on stream
printf(fmt, args)	Formatted standard output
sprintf(buf, fmt, args)	Formatted output to a string
vfprintf(stream, fmt, va_ptr)	Formatted output on stream
vprintf(fmt, va_ptr)	Formatted standard output
vsprintf(buf, fmt, va_ptr)	Formatted output to a string
fscanf(stream, fmt, args)	Formatted input from stream
scanf(fmt, args)	Formatted standard input
sscanf(buf, fmt, va_ptr)	Formatted input from a string
vfscanf(stream, fmt, va_ptr)	Formatted input from stream
vscanf(fmt, args)	Formatted standard input
vsscanf(buf, fmt, va_ptr)	Formatted input from a string
feof(stream)	True if stream at EOF
ferror(stream)	True if error on stream
clrerr(stream)	Reset error status on stream
fileno(stream)	Return fd from stream
remove(name)	Remove (delete) file

STRING HANDLING

atoi(s)	Convert ASCII decimal to integer
atol(s)	Convert ASCII decimal to long integer
atof(s)	Convert ASCII decimal to float
xtoi(s)	Convert ASCII hexadecimal to integer
memchr(s, c, n)	Find char in memory block
memcmp(s1, s2, n)	Compare n bytes of memory
memcpy(s1, s2, n)	Copy n bytes from s2 to s1
memmove(s1, s2, n)	Copy n bytes from s2 to s1
memset(s, c, n)	Set n bytes at s to c
strcat(s1, s2)	Append string 2 to string 1
strncat(s1, s2, n)	Append at most n chars to string 1
strcmp(s1, s2)	Compare strings
strncmp(s1, s2, n)	Compare n bytes of strings
strcpy(s1, s2)	Copy s2 to s1
strncpy(s1, s2, n)	Copy at most n bytes of s2
strerror(errno)	Map errno to an error message string
strlen(s)	Length of string
strchr(s, c)	Find char in string
strrchr(s, c)	Find rightmost char in string
strspn(s1, s2)	Length of s1 composed of chars from s2
strcspn(s1, s2)	Length of s2 composed of chars not from s2
strstr(s1, s2)	Locate the first occurrence of s2 in s1

LOW LEVEL I/O

open(name, mode)	Open a file
close(fd)	Close a file
creat(name)	Create a file
dup(fd)	Duplicate file descriptor
lseek(fd, offs, wh)	Random access positioning
read(fd, buf, cnt)	Read from file
rename(name1, name2)	Rename file
unlink(name)	Remove file from directory
write(fd, buf, cnt)	Write to file
isatty(fd)	True if fd refers to tty-like device
stat(name, buf)	Get information about a file
chmod(name, mode)	Set file attributes

CHARACTER TESTING

isalpha(c)	True if c is a letter
isupper(c)	Upper case letter
islower(c)	Lower case letter
isdigit(c)	Digit
isalnum(c)	Alphanumeric character
isspace(c)	Space, tab, newline, return or formfeed
ispunct(c)	Punctuation character
isprint(c)	Printable character
isgraph(c)	Printable non-space character
iscntrl(c)	Control character
isascii(c)	Ascii character (0-127)

FLOATING POINT

cos(f)	Cosine function
sin(f)	Sine function
tan(f)	Tangent function
acos(f)	Arc cosine function
asin(f)	Arc sine function
atan(f)	Arc tangent function
exp(f)	Exponential of f
log(f)	Natural log of f
log10(f)	Base 10 log of f
pow(x, y)	X to the y'th power
sqrt(f)	Square root
fabs(f)	Floating absolute value
ceil(f)	Smallest integral value $\geq f$
floor(f)	Largest integral value $\leq f$
sinh(f)	Hyperbolic sine
cosh(f)	Hyperbolic cosine
tanh(f)	Hyperbolic tangent
frexp(y, p)	Split into mantissa and exponent
ldexp(y, i)	Load new exponent

CONSOLE I/O

getch()	Get single character
getche()	Get single character with echo
putch(c)	Put single character
ungetch(c)	Push character back
kbhit()	Test for key pressed
cgets(s)	Get line from console
cputs(s)	Put string to console

DATE AND TIME FUNCTIONS

time(p)	Get current date/time
gmtime(p)	Get broken down Universal time
localtime(p)	Get broken down local time
asctime(t)	Convert broken down time to ascii
ctime(p)	Convert time to ascii

MISCELLANEOUS

execl(name, args)	Execute another program
execv(name, argp)	Execute another program
spawnl(name, arg, ...)	Execute a subprogram
spawnv(name, argp)	Execute a subprogram
system(s)	Execute system command
atexit(func)	Install func to be executed on termination
exit(status)	Terminate execution
_exit(status)	Terminate execution immediately
getuid()	Get user id (CP/M)
setuid(uid)	Set user id (CP/M)
chdir(s)	Change directory (MS-DOS)
mkdir(s)	Create directory (MS-DOS)
rmdir(s)	Remove directory (MS-DOS)
getcwd(drive)	Get current working directory (MS-DOS)
signal(sig, func)	Set trap for interrupt condition
brk(addr)	Set memory allocation
sbrk(incr)	Adjust memory allocation
malloc(cnt)	Dynamic memory allocation
free(ptr)	Dynamic memory release
realloc(ptr, cnt)	Dynamic memory reallocation
calloc(cnt, size)	Dynamic memory allocation zeroed
perror(s)	Print error message
qsort(base, nel, width, func)	Quick sort
srand(seed)	Initialize random number generator
rand()	Get next random number
setjmp(buf)	Setup for non-local goto
longjmp(buf, val)	Non-local goto
_getargs(buf, name)	Wild card expansion and i/o redirection
inp(port)	Read port
outp(port, data)	Write data to port
bdos(func, val)	Perform bdos call (CP/M)
msdos(func, val, val, ...)	Perform msdos call
msdoscx(func, val, val, ...)	Alternate msdos call
intdos(ip, op)	Execute DOS interrupt

<code>intdosx(ip, op, sp)</code>	Execute DOS interrupt
<code>segread(sp)</code>	Get segment register values
<code>int86(int, ip, op)</code>	Execute software interrupt
<code>int86x(int, ip, op, sp)</code>	Execute software interrupt
<code>bios(n, c)</code>	Call bios entry (CP/M)
<code>ei()</code>	Enable interrupts
<code>di()</code>	Disable interrupts
<code>set_vector(vec, func)</code>	Set an interrupt vector
<code>assert(e)</code>	Run time assertion
<code>getenv(s)</code>	Get environment string (MS-DOS)

ACOS, ASIN, ATAN, ATAN2

SYNOPSIS

```
#include <math.h>

double    acos(double f)
double    asin(double f)
double    atan(double f)

double    atan2(double x, double y)
```

DESCRIPTION

These functions are the converse of the trigonometric functions `cos`, `sin` and `tan`. `Acos` and `asin` are undefined for arguments whose absolute value is greater than 1.0. The returned value is in radians, and always in the range $-\pi/2$ to $+\pi/2$, except for `cos()`, which returns a value in the range 0 to π . `Atan2()` returns the inverse tan of x/y but uses the signs of its arguments to return a value in the range $-\pi$ to $+\pi$.

SEE ALSO

`sin`, `cos`, `tan`

ATEXIT

SYNOPSIS

```
#include <stdlib.h>

int atexit(void (*func)(void));
```

DESCRIPTION

The `atexit()` function registers the function pointed to by `func`, to be called without arguments at normal program termination. `atexit()` returns zero if the registration succeeds, nonzero if it fails. On program termination, all functions registered by `atexit()` are called, in the reverse order of their registration.

SEE ALSO

`exit`

ASCTIME

SYNOPSIS

```
#include <time.h>

char * asctime(time_t t)
```

DESCRIPTION

`Asctime()` takes the broken down time pointed to by its argument, and returns a 26 character string describing the current date and time in the format

```
Sun Sep 16 01:03:52 1973\n\0
```

Note the newline at the end of the string. The width of each field in the string is fixed.

SEE ALSO

`ctime, time, gmtime, localtime`

ASSERT

SYNOPSIS

```
#include <assert.h>

void      assert(int e)
```

DESCRIPTION

This macro is used for debugging purposes; the basic method of usage is to place assertions liberally throughout your code at points where correct operation of the code depends upon certain conditions being true initially. An `assert()` may be used to ensure at run time that that assumption holds. For example, the following statement asserts that the pointer `tp` is non-null:

```
assert(tp);
```

If at run time the expression evaluates to false, the program will abort with a message identifying the source file and line number of the assertion, and the expression used as an argument to it. A fuller discussion of the uses of `assert` is impossible in limited space, but it is closely linked to methods of proving program correctness.

ATOF, ATOI, ATOL

SYNOPSIS

```
#include <math.h>

double    atof(char * s)
int atoi(char * s)

#include <stdlib.h>

long      atol(char * s)
```

DESCRIPTION

These routines convert a decimal number in the argument string `s` into a double float, integer or long integer respectively. Leading blanks are skipped over. In the case of `atof()`, the number may be in scientific notation.

BDOS (CP/M only)

SYNOPSIS

```
#include <cpm.h>

char      bdos(int func, int arg)

short bdoshl(int func, int arg) (CP/M-80 only)
```

DESCRIPTION

Bdos() calls the CP/M BDOS with func in register C (CL for CP/M-86) and arg in register DE (DX). The return value is the byte returned by the BDOS in register A (AX). Bdoshl() is the same, except that the return value is the value returned by the BDOS in HL. Constant values for the various BDOS function values are defined in cpm.h.

These functions should be avoided except in programs which are not intended to be used on an operating system other than CP/M. The standard I/O routines are to be preferred, since they are portable.

SEE ALSO

bios, msdos

BIOS (CP/M only)

SYNOPSIS

```
#includ  <cpm.h>

char bios(int n, int a1, int a2)
```

DESCRIPTION

This function will call the n'th bios entry point (cold boot = 0, warm boot = 1, etc.) with register BC (CX) set to the argument a1 and DE (DX) set to the argument a2. The return value is the contents of register A (AX) after the bios call. On CP/M-86, bdos function 50 is used to perform the bios call. This function should not be used unless unavoidable, since it is highly non-portable. There is even no guarantee of portability of bios calls between differing CP/M systems.

SEE ALSO

bdos

CALLOC

SYNOPSIS

```
#include <stdlib.h>

char * calloc(size_t cnt, size_t size)
```

DESCRIPTION

Calloc() attempts to obtain a contiguous block of dynamic memory which will hold cnt objects, each of length size. The block is filled with zeroes. A pointer to the block is returned, or 0 if the memory could not be allocated.

SEE ALSO

brk, sbrk, malloc, free

CGETS, CPUTS

SYNOPSIS

```
#include <conio.h>

char *  cgets(char * s)

void    cputs(char * s)
```

DESCRIPTION

Cputs() will read one line of input from the console into the buffer passed as an argument. It does so by repeated calls to getche(). Cputs() writes its argument string to the console, outputting carriage returns before each newline in the string. It calls putch() repeatedly.

SEE ALSO

getch, getche, putch

CHDIR

SYNOPSIS

```
#include <sys.h>

int chdir(char * s)
```

DESCRIPTION

This function is available only under MS-DOS. It changes the current working directory to the path name supplied as argument. This path name may be absolute, as in A:\FRED, or relative, as in ..\SOURCES. A return value of -1 indicates that the requested change could not be performed.

SEE ALSO

mkdir, rmdir, getcwd

CHMOD

SYNOPSIS

```
#include <stat.h>

int chmod(char * name, int )
char *   name;
int mode;
```

DESCRIPTION

This function changes the file attributes (or modes) of the named file. The argument name may be any valid file name. The mode argument may include all bits defined in stat.h except those relating to the type of the file, e.g. S_IFDIR. Note however that not all bits may be changed under all operating systems, e.g. neither DOS nor CP/M permit a file to be made unreadable, thus even if mode does not include S_IREAD the file will still be readable (and stat() will still return S_IREAD in flags).

SEE ALSO

stat, creat

CLOSE

SYNOPSIS

```
#include <unistd.h>

int close(int fd)
```

DESCRIPTION

This routine closes the file associated with the file descriptor `fd`, which will have been previously obtained from a call to `open()`. `Close()` returns 0 for a successful close, or -1 otherwise.

SEE ALSO

`open`, `read`, `write`, `seek`

CLRERR, CLREOF

SYNOPSIS

```
#include <stdio.h>

void clrerr(FILE * stream)
void clreof(FILE * stream)
```

DESCRIPTION

These are macros, defined in `stdio.h`, which reset the error and end of file flags respectively for the specified stream. They should be used with care; the major valid use is for clearing an EOF status on input from a terminal-like device, where it may be valid to continue to read after having seen an end-of-file indication.

SEE ALSO

`fopen`, `fclose`

COS

SYNOPSIS

```
#include <math.h>

double    cos(double f)
```

DESCRIPTION

This function yields the cosine of its argument.

SEE ALSO

sin, tan, asin, acos, atan

COSH, SINH, TANH

SYNOPSIS

```
#include <math.h>

double    cosh(double f)
double    sinh(double f)
double    tanh(double f)
```

DESCRIPTION

These functions implement the hyperbolic trig functions.

CREAT

SYNOPSIS

```
#include <stat.h>

int creat(char * name, int mode)
```

DESCRIPTION

This routine attempts to create the file named by name. If the file exists and is writeable, it will be removed and re-created. The return value is -1 if the create failed, or a small non-negative number if it succeeded. This number is a valuable token which must be used to write to or close the file subsequently. Mode is used to initialize the attributes of the created file. The allowable bits are the same as for chmod(), but for Unix compatibility it is recommended that a mode of 0666 or 0600 be used. Under CP/M the mode is ignored - the only way to set a file's attributes is via the chmod() function.

SEE ALSO

open, close, read, write, seek, stat, chmod

CTIME

SYNOPSIS

```
#include <time.h>

char *    ctime(time_t t)
```

DESCRIPTION

Ctime() converts the time in seconds pointed to by its argument to a string of the same form as described for asctime. Thus the following program prints the current time and date:

```
#include <time.h>

main()
{
    time_t    t;

    time(&t);
    printf("%s", ctime(&t));
}
```

SEE ALSO

gmtime, localtime, asctime, time

DIV, LDIV

SYNOPSIS

```
#include <stdlib.h>

div_t      div(int numer, int denom)
ldiv_t     ldiv(long numer, long denom)
```

DESCRIPTION

The `div()` function computes the quotient and remainder of the division of `numer` by `denom`. The `div()` function returns a structure of type `div_t`, containing both the quotient and remainder. `ldiv()` is similar to `div()` except it takes arguments of type `long` and returns a structure of type `ldiv_t`. The types `div_t` and `ldiv_t` are defined in `<stdlib.h>` as follows:

```
typedef struct {
    intquot, rem;
} div_t;

typedef struct {
    longquot, rem;
} ldiv_t;
```

DI, EI

SYNOPSIS

```
void ei(void);  
void di(void);
```

DESCRIPTION

Ei() and di() enable and disable interrupts respectively.

DUP

SYNOPSIS

```
#include <unixio.h>  
  
int dup(int fd)
```

DESCRIPTION

Given a file descriptor, such as returned by open(), this routine will return another file descriptor which will refer to the same open file. -1 is returned if the fd argument is a bad descriptor or does not refer to an open file.

SEE ALSO

open, close, creat, read, write

EXECL, EXECV

SYNOPSIS

```
#include <sys.h>

int execl(char * name, pname, ...)
int execv(char * name, ppname)
```

DESCRIPTION

Execl() and execv() load and execute the program specified by the string name. Execl() takes the arguments for the program from the zero-terminated list of string arguments. Execv() is passed a pointer to an array of strings. The array must be zero-terminated. If the named program is found and can be read, the call does not return. Thus any return from these routines may be treated as an error.

SEE ALSO

spawnl, spawnv, system

EXIT

SYNOPSIS

```
#include <stdlib.h>

void      exit(int status)
```

DESCRIPTION

This call will close all open files and exit from the program. On CP/M, this means a return to CCP level. Status will be stored in a known place for examination by other programs. This is only useful if the program executing was actually invoked by another program which is trapping warm boots. The status value will be stored on CP/M at 80H. This call will never return.

SEE ALSO

atexit

_EXIT

SYNOPSIS

```
#include <stdlib.h>
void      _exit(int status)
```

DESCRIPTION

This function will cause an immediate exit from the program, without the normal flushing of stdio buffers that is performed by `exit()`.

SEE ALSO

`exit`

EXP, LOG, LOG10, POW

SYNOPSIS

```
#include <math.h>

double    exp(double f)
double    log(double f)
double    log10(double f)
double    pow(double x, y)
```

DESCRIPTION

`Exp()` returns the exponential function of its argument, `log()` the natural logarithm of `f`, and `log10()` the logarithm to base 10. `Pow()` returns the value of `x` raised to the `y`'th power.

FABS, CEIL, FLOOR

SYNOPSIS

```
#include <math.h>

double    fabs(double f)
double    ceil(double f)
double    floor(double f)
```

DESCRIPTION

These routines return respectively the absolute value of *f*, the smallest integral value not less than *f*, and the largest integral value not greater than *f*.

FCLOSE

SYNOPSIS

```
#include <stdio.h>

int fclose(FILE * stream)
```

DESCRIPTION

This routine closes the specified i/o stream. Stream should be a token returned by a previous call to `fopen()`. NULL is returned on a successful close, EOF otherwise.

SEE ALSO

`fopen`, `fread`, `fwrite`

FEOF, FERROR

SYNOPSIS

```
#include <stdio.h>

feof(FILE * stream)
ferror(FILE * stream)
```

DESCRIPTION

These macros test the status of the EOF and ERROR bits respectively for the specified stream. Each will be true if the corresponding flag is set. The macros are defined in `stdio.h`. Stream must be a token returned by a previous `fopen()` call.

SEE ALSO

`fopen`, `fclose`

FFLUSH

SYNOPSIS

```
#include <stdio.h>

int fflush(FILE * stream)
```

DESCRIPTION

`fflush()` will output to the disk file or other device currently open on the specified stream the contents of the associated buffer. This is typically used for flushing buffered standard output in interactive applications.

SEE ALSO

`fopen`, `fclose`

FGETC

SYNOPSIS

```
#include <stdio.h>

int fgetc(FILE * stream)
```

DESCRIPTION

Fgetc() returns the next character from the input stream. If end-of-file is encountered EOF will be returned instead. It is for this reason that the function is declared as int. The integer EOF is not a valid byte, thus end-of-file is distinguishable from reading a byte of all 1 bits from the file. Fgetc() is the non-macro version of getc().

SEE ALSO

fopen, fclose, fputc, getc, putc

FGETS

SYNOPSIS

```
#include <stdio.h>

char * fgets(char * s, size_t n, char * stream)
```

DESCRIPTION

Fgets() places in the buffer s up to n-1 characters from the input stream. If a newline is seen in the input before the correct number of characters is read, then fgets() will return immediately. The newline will be left in the buffer. The buffer will be null terminated in any case. A successful fgets() will return its first argument; NULL is returned on end-of-file or error.

FILENO

SYNOPSIS

```
fileno(FILE * stream)
```

DESCRIPTION

Fileno() is a macro from stdio.h which yields the file descriptor associated with stream. It is mainly used when it is desired to perform some low-level operation on a file opened as a stdio stream.

SEE ALSO

fopen, fclose, open, close

FOPEN

SYNOPSIS

```
#include <stdio.h>

FILE * fopen(char * name, char * mode);
```

DESCRIPTION

DESCRIPTION

Fopen() attempts to open file for reading or writing (or both) according to the mode string supplied. The mode string is interpreted as follows:

- r The file is opened for reading if it exists. If the file does not exist the call fails.
- r+ If the file exists it is opened for reading and writing. If the file does not already exist the call fails.
- w The file is created if it does not exist, or truncated if it does. It is then opened for writing.
- w+ The file is created if it does not already exist, or truncated if it does. The file is opened for reading and writing.
- a The file is created if it does not already exist, and opened for writing. All writes will be dynamically forced to the end of file, thus this mode is known as append mode.
- a+ The file is created if it does not already exist, and opened for reading and writing. All writes to the file will be dynamically forced to the end of the file, i.e. while any portion of the file may be read, all writes

will take place at the end of the file and will not overwrite any existing data. Calling `fseek()` in an attempt to write at any other place in the file will not be effective.

The "b" modifier may be appended to any of the above modes, e.g. "r+b" or "rb+" are equivalent. Adding the "b" modifier will cause the file to be opened in binary rather than ASCII mode. Opening in ASCII mode ensures that text files are read in a manner compatible with the Unix-derived conventions for C programs, i.e. that text files contain lines delimited by newline characters. The special treatment of read or written characters varies with the operating system, but includes some or all of the following:

NEWLINE (LINE FEED)

Converted to carriage return, line feed on output.

RETURN

Ignored on input, inserted before NEWLINE on output.

CTRL-Z

Signals EOF on input, appended on `fclose` on output if necessary on CP/M.

Opening a file in binary mode will allow each character to be read just as written, but because the exact size of a file is not known to CP/M, the file may contain more bytes than were written to it. See `open()` for a description of what constitutes a file name.

When using one of the read/write modes (with a '+' character in the string), although they permits reading and writing on the same stream, it is not possible to arbitrarily mix input and output calls to the same stream. At any given time a stream opened with a "+" mode will be in either an input or output state. The state may only be changed when the associated buffer is empty, which is only guaranteed immediately after a call to `fflush()` or one of the file positioning functions `fseek()` or `rewind()`. The buffer will also be empty after encountering EOF while reading a binary stream, but it is recommended that an explicit call to `fflush()` be used to ensure this situation. Thus after reading from a stream you should call `fflush()` or `fseek()` before attempting to write on that stream, and vice versa.

SEE ALSO

`fclose`, `fgetc`, `fputc`, `freopen`

FPRINTF

SYNOPSIS

```
#include <stdio.h>

fprintf(FILE * stream, char * fmt, ...);
vfprintf(FILE * stream, va_list va_arg);
```

DESCRIPTION

Fprintf() performs formatted printing on the specified stream. Refer to printf() for the details of the available formats. Vfprintf() is similar to fprintf() but takes a variable argument list pointer rather than a list of arguments. See the description of va_start() for more information on variable argument lists.

SEE ALSO

printf, fscanf, sscanf

FPUTC

SYNOPSIS

```
#include <stdio.h>

int fputc(int c, FILE * stream)
```

DESCRIPTION

The character c is written to the supplied stream. This is the non-macro version of putc(). The character is returned if it was successfully written, EOF is returned otherwise. Note that "written to the stream" may mean only placing the character in the buffer associated with the stream.

SEE ALSO

putc, fgetc, fopen, fflush

FPUTS

SYNOPSIS

```
#include <stdio.h>

int fputs(char * s, FILE * stream)
```

DESCRIPTION

The null-terminated string *s* is written to the *stream*. No newline is appended (cf. `puts()`). The error return is EOF.

SEE ALSO

`puts, fgets, fopen, fclose`

FREAD

SYNOPSIS

```
#include <stdio.h>

int fread(void * buf, size_t size, size_t cnt,
          FILE * stream)
```

DESCRIPTION

Up to *cnt* objects, each of length *size*, are read into memory at *buf* from the *stream*. The return value is the number of objects read. If none is read, 0 will be returned. Note that a return value less than *cnt*, but greater than 0, may not represent an error (cf. `fwrite()`). No word alignment in the stream is assumed or necessary. The read is done via successive `getc()`'s.

SEE ALSO

`fwrite, fopen, fclose, getc`

FREE

SYNOPSIS

```
#include <stdlib.h>

void      free(void * ptr)
```

DESCRIPTION

Free() deallocates the block of memory at ptr, which must have been obtained from a call to malloc() or calloc().

SEE ALSO

malloc, calloc

FREOPEN

SYNOPSIS

```
#include <stdio.h>

FILE * freopen(char * name, char * mode, FILE * stream)
```

DESCRIPTION

Freopen() closes the given stream (if open) then re-opens the stream attached to the file described by name. The mode of opening is given by mode. It either returns the stream argument, if successful, or NULL if not. See fopen() for more information.

SEE ALSO

fopen, fclose

FREXP, LDEXP

SYNOPSIS

```
#include <math.h>

double    frexp(double f, int * p)

double    ldexp(double f, int i)
```

DESCRIPTION

Frexp() breaks a floating point number into a normalized fraction and an integral power of 2. The integer is stored into the int object pointed to by p. Its return value x is in the interval [0.5, 1.0) or zero, and f equals x times 2 raised to the power stored in *p. If f is zero, both parts of the result are zero. Ldexp() performs the reverse operation; the integer i is added to the exponent of the floating point f and the resultant value returned.

FSCANF

SYNOPSIS

```
#include <stdio.h>

int fscanf(FILE * stream, char * fmt, ...)
```

DESCRIPTION

This routine performs formatted input from the specified stream. See scanf() for a full description of the behaviour of the routine. Vscanf() is similar to fscanf() but takes a variable argument list pointer rather than a list of arguments. See the description of va_start() for more information on variable argument lists.

SEE ALSO

scanf, sscanf, fopen, fclose

FSEEK

SYNOPSIS

```
#include <stdio.h>

int fseek(FILE * stream, long offs, int wh)
```

DESCRIPTION

Fseek() positions the "file pointer" (i.e. a pointer to the next character to be read or written) of the specified stream as follows:

w_h	r_e_s_u_l_t_a_n_t	l_o_c_a_t_i_o_n
0	offs	
1	offs+previous location	
2	offs+length of file	

It should be noted that offs is a signed value. Thus the 3 allowed modes give positioning relative to the beginning of the file, the current file pointer and the end of the file respectively. EOF is returned if the positioning request could not be satisfied. Note however that positioning beyond the end of the file is legal, but will result in an EOF indication if an attempt is made to read data there. It is quite in order to write data beyond the previous end of file. Fseek() correctly accounts for any buffered data.

SEE ALSO

lseek, fopen, fclose

FTELL

SYNOPSIS

```
#include <stdio.h>

long      ftell(FILE * stream)
```

DESCRIPTION

This function returns the current position of the conceptual read/write pointer associated with stream. This is the position relative to the beginning of the file of the next byte to be read from or written to the file.

SEE ALSO

fseek

FWRITE

SYNOPSIS

```
#include <stdio.h>

int fwrite(void * buf, size_t size, size_t cnt,
           FILE * stream)
```

DESCRIPTION

Cnt objects of length size bytes will be written from memory at buf, to the specified stream. The number of whole objects written will be returned, or 0 if none could be written. Any return value not equal to cnt should be treated as an error (cf. fread()).

SEE ALSO

fread, fopen, fclose

_GETARGS

SYNOPSIS

```
#include <sys.h>

char ** _getargs(char * buf, char * name)
extern int _argc_;
```

DESCRIPTION

This routine performs I/O redirection (CP/M only) and wild card expansion. Under MS-DOS I/O redirection is performed by the operating system. It is called from startup code to operate on the command line if the -R option is used to the C command, but may also be called by user-written code. If the buf argument is null, it will read lines of text from standard input. If the standard input is a terminal (usually the console) the name argument will be written to the standard error stream as a prompt. If the buf argument is not null, it will be used as the source of the string to be processed. The returned value is a pointer to an array of strings, exactly as would be pointed to by the argv argument to the main() function. The number of strings in the array may be obtained from the global _argc_. For example, a typical use of this function would be:

```
#include <sys.h>

main(argc, argv)
char ** argv;
{
    extern char ** _getargs();
    extern int _argc_;

    if(argc == 1) { /* no arguments */
        argv = _getargs(0, "myname");
        argc = _argc_;
    }
    .
    .
    .
}
```

There will be one string in the array for each word in the buffer processed. Quotes, either single (') or double (") may be used to include white space in "words". If any wild card characters (? or *) appear in a non-quoted word, it will be expanded into a string of words, one for each file matching the word. The usual CP/M conventions are followed for this expansion. On CP/M any occurrence of the redirection characters > and < outside quotes will be handled in the following manner:

```
> name
    will cause standard output to be redirected to the file
    name.
```

< name
will cause standard input to be redirected from the file name.

>> name
will cause standard output to append to file name.

White space is optional between the > or < character and the file name, however it is an error for a redirection character not to be followed by a file name. It is also an error if a file cannot be opened for input or created for output. An append redirection (>>) will create the file if it does not exist. If the source of text to be processed is standard input, several lines may be supplied by ending each line (except the last) with a backslash (\). This serves as a continuation character. Note that the newline following the backslash is ignored, and not treated as white space.

GETC

SYNOPSIS

```
#include <stdio.h>

int getc(FILE * stream)
FILE * stream;
```

DESCRIPTION

One character is read from the specified stream and returned. EOF will be returned on end-of-file or error. This is the macro version of fgetc(), and is defined in stdio.h.

GETCH, GETCHE, UNGETCH, PUTCH

SYNOPSIS

```
#include <conio.h>

char    getch(void)
char    getche(void)
void    putchar(int c)
```

DESCRIPTION

Getch() reads a single character from the console keyboard and returns it without echoing. Getche() is similar but does echo the character typed. Ungetch() will push back one character such that the next call to getch() or getche() will return that character. Putch() outputs the character c to the console screen, prepending a carriage return if the character is a new-line.

SEE ALSO

cgets, cputs

GETCHAR

SYNOPSIS

```
#include <stdio.h>

int getchar(void)
```

DESCRIPTION

Getchar() is a getc(stdin) operation. It is a macro defined in stdio.h. Note that under normal circumstances getchar() will NOT return unless a carriage return has been typed on the console. To get a single character immediately from the console, use the routine getch().

SEE ALSO

getc, fgetc, freopen, fclose

GETCWD (MS-DOS only)

SYNOPSIS

```
#include <sys.h>

char *    getcwd(int drive)
```

DESCRIPTION

Getcwd() returns the path name of the current working directory on the specified drive, where drive == 0 represents the current drive, drive == 1 represents A:, drive == 2 represents B: etc. The return value is a pointer to a static area of memory which will be overwritten on the next call to getcwd().

SEE ALSO

chdir

GETENV

SYNOPSIS

```
#include <stdlib.h>

char *    getenv(char * s)
extern char ** environ;
```

DESCRIPTION

Getenv() will search the vector of environment strings for one matching the argument supplied, and return the value part of that environment string. For example, if the environment contains the string

```
COMSPEC=A:\COMMAND.COM
```

then getenv("COMSPEC") will return A:\COMMAND.COM. The global variable environ is a pointer to an array of pointers to environment strings, terminated by a null pointer. This array is initialized at startup time under MS-DOS from the environment pointer supplied when the program was executed. Under CP/M no such environment is supplied, so the first call to getenv() will attempt to open a file in the current user number on the current drive called ENVIRON. This file should contain definitions for any environment variables desired to be accessible to the program, e.g.

```
HITECH=0:C:
```

Each variable definition should be on a separate line, consisting of the variable name (conventionally all in upper case) followed without intervening white space by

an equal sign ('=') then the value to be assigned to that variable.

GETS

SYNOPSIS

```
#include <stdio.h>

char * gets(char * s)
```

DESCRIPTION

Gets() reads a line from standard input into the buffer at s, deleting the newline (cf. fgets()). The buffer is null terminated. It returns its argument, or NULL on end-of-file.

SEE ALSO

fgets, freopen

GETUID (CP/M only)

SYNOPSIS

```
#include <sys.h>

int getuid(void)
```

DESCRIPTION

Getuid() returns the current user number. On CP/M, the current user number determines the user number associated with an opened or created file, unless overridden by an explicit user number prefix in the file name.

SEE ALSO

setuid, open

GETW

SYNOPSIS

```
#include <stdio.h>

int getw(FILE * stream)
```

DESCRIPTION

Getw() returns one word (16 bits for the Z80 and 8086) from the nominated stream. EOF is returned on end-of-file, but since this is a perfectly good word, the feof() macro should be used for testing for end-of-file. When reading the word, no special alignment in the file is necessary, as the read is done by two consecutive getc()'s. The byte ordering is however undefined. The word read should in general have been written by putw().

SEE ALSO

putw, getc, fopen, fclose

GMTIME, LOCALTIME

SYNOPSIS

```
#include <time.h>

struct tm *      gmtime(time_t * t)
struct tm *      localtime(time_t * t)
```

DESCRIPTION

These functions convert the time pointed to by t which is in seconds since 00:00:00 on Jan 1, 1970, into a broken down time stored in a structure as defined in time.h. Gmtime() performs a straight conversion, while localtime() takes into account the contents of the global integer time_zone. This should contain the number of minutes that the local time zone is WESTWARD of Greenwich. Since there is no way under MS-DOS of actually pre-determining this value, by default localtime() will return the same result as gmtime().

SEE ALSO

ctime, asctime, time

INP, OUTP

SYNOPSIS

```
char inp(unsigned port)

void outp(unsigned, unsigned data)
```

DESCRIPTION

These routines read and write bytes to and from I/O ports. Inp() returns the data byte read from the specified port, and outp() outputs the data byte to the specified port.

INT86, INT86X, INTDOS, INTDOSX

SYNOPSIS

```
#include <dos.h>

int int86(int intno, union REGS * inregs,
          union REGS * outregs)
int int86x(int intno, union REGS inregs,
          union REGS outregs, struct SREGS * segregs)
int intdos(union REGS * inregs, union REGS * outregs)
int intdosx(union REGS * inregs, union REGS * outregs,
            struct SREGS * segregs)
```

DESCRIPTION

These functions allow calling of software interrupts from C programs. Int86() and int86x() execute the software interrupt specified by intno while intdos() and intdosx() execute interrupt 21(hex), which is the MS-DOS system call interrupt. The inregs pointer should point to a union containing values for each of the general purpose registers to be set when executing the interrupt, and the values of the registers on return are copied into the union pointed to by outregs. The x versions of the calls also take a pointer to a union defining the segment register values to be set on execution of the interrupt, though only ES and DS are actually set from this structure.

SEE ALSO

segread

ISALNUM, ISALPHA, ISDIGIT, ISLOWER et. al.

SYNOPSIS

```
#include <ctype.h>

isalnum(char c)
isalpha(char c)
isascii(char c)
iscntrl(char c)
isdigit(char c)
islower(char c)
isprint(char c)
isgraph(char c)
ispunct(char c)
isspace(char c)
isupper(char c)
char c;
```

DESCRIPTION

These macros, defined in `ctype.h`, test the supplied character for membership in one of several overlapping groups of characters. Note that all except `isascii` are defined for `c` iff `isascii(c)` is true.

<code>isalnum(c)</code>	<code>c</code> is alphanumeric
<code>isalpha(c)</code>	<code>c</code> is in A-Z or a-z
<code>isascii(c)</code>	<code>c</code> is a 7 bit ascii character
<code>iscntrl(c)</code>	<code>c</code> is a control character
<code>isdigit(c)</code>	<code>c</code> is a decimal digit
<code>islower(c)</code>	<code>c</code> is in a-z
<code>isprint(c)</code>	<code>c</code> is a printing char
<code>isgraph(c)</code>	<code>c</code> is a non-space printable character
<code>ispunct(c)</code>	<code>c</code> is not alphanumeric
<code>isspace(c)</code>	<code>c</code> is a space, tab or newline
<code>isupper(c)</code>	<code>c</code> is in A-Z
<code>isxdigit(c)</code>	<code>c</code> is in 0-9 or a-f or A-F

SEE ALSO

`toupper`, `tolower`, `toascii`

ISATTY

SYNOPSIS

```
#include <unixio.h>

int isatty(int fd)
```

DESCRIPTION

This tests the type of the file associated with `fd`. It returns true if the file is attached to a tty-like device. This would normally be used for testing if standard input is coming from a file or the console. For testing STDIO streams, use `isatty(fileno(stream))`.

KBHIT

SYNOPSIS

```
#include <conio.h>

int kbhit(void)
```

DESCRIPTION

This function returns 1 if a character has been pressed on the console keyboard, 0 otherwise. Normally the character would then be read via `getch()`.

SEE ALSO

`getch`, `getche`

LONGJMP

SYNOPSIS

```
#include <setjmp.h>

void      longjmp(jmp_buf buf, int val)
```

DESCRIPTION

Longjmp(), in conjunction with setjmp(), provides a mechanism for non-local gotos. To use this facility, setjmp() should be called with a jmp_buf argument in some outer level function. The call from setjmp() will return 0. To return to this level of execution, longjmp() may be called with the same jmp_buf argument from an inner level of execution. Note however that the function which called setjmp() must still be active when longjmp() is called. Breach of this rule will cause disaster, due to the use of a stack containing invalid data. The val argument to longjmp() will be the value apparently returned from the setjmp(). This should normally be non-zero, to distinguish it from the genuine setjmp() call. For example:

```
#include <setjmp.h>

static jmp_buf  jb_err;

main()
{
    if(setjmp(jb_err)) {
        printf("An error occurred");
        exit(1);
    }
    a_func();
}

a_func()
{
    if(do_whatever() != 0)
        longjmp(jb_err, 1);
    if(do_something_else() != 0)
        longjmp(jb_err, 2);
}
```

The calls to longjmp() above will never return; rather the call to setjmp() will appear to return, but with a return value equal to the argument supplied to longjmp().

SEE ALSO

setjmp

LSEEK

SYNOPSIS

```
#include <unixio.h>

long lseek(int fd, long offs, int wh)
```

DESCRIPTION

This function operates in an analogous manner to `fseek()`, however it does so on unbuffered low-level i/o file descriptors, rather than on `STDIO` streams. It also returns the resulting pointer location. Thus `lseek(fd, 0L, 1)` returns the current pointer location without moving it. `-1` is returned on error.

SEE ALSO

`open`, `close`, `read`, `write`

MALLOC

SYNOPSIS

```
#include <stdlib.h>

void * malloc(size_t cnt)
```

DESCRIPTION

`Malloc()` attempts to allocate `cnt` bytes of memory from the "heap", the dynamic memory allocation area. If successful, it returns a pointer to the block, otherwise `0` is returned. The memory so allocated may be freed with `free()`, or changed in size via `realloc()`. `Malloc()` calls `sbrk()` to obtain memory, and is in turn called by `calloc()`. `Malloc()` does not clear the memory it obtains.

SEE ALSO

`calloc`, `free`, `realloc`

MEMSET, MEMCPY, MEMCMP, MEMMOVE

SYNOPSIS

```
#include <string.h>

void      memset(void s, char c, size_t n)
void *    memcpy(void * d, void * s, size_t n)
int memcmp(void * s1, void * s2, size_t n)
void *    memmove(void * s1, void * s2, size_t n)
void *    memchr(void * s, int c, size_t n)
```

DESCRIPTION

Memset() initializes n bytes of memory starting at the location pointed to by s with the character c. Memcpy() copies n bytes of memory starting from the location pointed to by s to the block of memory pointed to by d. The result of copying overlapping blocks is undefined. Memcmp() compares two blocks of memory, of length n, and returns a signed value similar to strcmp(). Unlike strcmp() the comparison does not stop on a null character. The ascii collating sequence is used for the comparison, but the effect of including non-ascii characters in the memory blocks on the sense of the return value is indeterminate. Memmove() is similar to memcpy() except copying of overlapping blocks is handled correctly. The memchr() function locates the first occurrence of c (converted to unsigned char) in the initial n characters of the object pointed to by s.

SEE ALSO

strcpy, strcmp, strchr

MKDIR, RMDIR

SYNOPSIS

```
#include <sys.h>

int mkdir(char * s)
int rmdir(char * s)
```

DESCRIPTION

These functions allow the creation (mkdir()) and deletion (rmdir()) of sub-directories under the MS-DOS operating system. The argument *s* may be an arbitrary pathname, and the return value will be -1 if the creation or removal was unsuccessful.

SEE ALSO

chdir

MSDOS, MSDOSCX

SYNOPSIS

```
#include <dos.h>

long      msdos(int ax, int dx, int cx,
               int bx, int si, int di)
long      msdoscx(int ax, int dx, int cx,
                 int bx, int si, int di)
```

DESCRIPTION

These functions allow direct access to MS-DOS system calls. The arguments will be placed in the registers implied by their names, while the return value will be the contents of AX and DX (for msdos()) or the contents of DX and CX (for msdoscx()). Only as many arguments as necessary need be supplied, e.g. if only AH and DX need have specified valued, then only 2 argument would be required. The following piece of code outputs a form-feed to the printer.

```
msdos(0x500, '\f');
```

Note that the system call number (in this case 5) must be multiplied by 0x100 since MS-DOS expects the call number in AH, the high byte of AX.

SEE ALSO

intdos, intdosx, int86, int86x

OPEN

SYNOPSIS

```
#include <unixio.h>

int open(char * name, int mode)
```

DESCRIPTION

Open() is the fundamental means of opening files for reading and writing. The file specified by name is sought, and if found is opened for reading, writing or both. Mode is encoded as follows:

Mode	Meaning
0	Open for reading only
1	Open for writing only
2	Open for both reading and writing

The file must already exist - if it does not, creat() should be used to create it. On a successful open, a file descriptor is returned. This is a non-negative integer which may be used to refer to the open file subsequently. If the open fails, -1 is returned. The syntax of a CP/M filename is:

```
[uid:][drive:]name.type
```

where uid is a decimal number 0 to 15, drive is a letter A to P or a to p, name is 1 to 8 characters and type is 0 to 3 characters. Though there are few inherent restrictions on the characters in the name and type, it is recommended that they be restricted to the alphanumerics and standard printing characters. Use of strange characters may cause problems in accessing and/or deleting the file.

One or both of uid: and drive: may be omitted; if both are supplied, the uid: must come first. Note that the [and] are meta-symbols only. Some examples are:

```
fred.dat
file.c
0:xyz.com
0:a:file1.p
a:file2.
```

If the uid: is omitted, the file will be sought with uid equal to the current user number, as returned by getuid(). If drive: is omitted, the file will be sought on the currently selected drive. The following special file names are recognized:

```
lst:      Accesses the list device - write only
pun:      Accesses the punch device - write only
rdr:      Accesses the reader device - read only
con:      Accesses the system console - read/write
```

File names may be in any case - they are converted to upper case during processing of the name.

MS-DOS filenames may be any valid MS-DOS 2.xx filename, e.g.

```
fred.nrk
A:\HITECH\STDIO.H
```

The special device names (e.g. CON, LST) are also recognized. These do not require (and should not have) a trailing colon.

SEE ALSO

close, fopen, fclose, read, write, creat

PERROR

SYNOPSIS

```
#include <stdio.h>

void      perror(char * s)
```

DESCRIPTION

This routine will print on the stderr stream the argument s, followed by a descriptive message detailing the last error returned from an open, close read or write call. Unfortunately CP/M does not provide definitive information relating to the error, except in the case of a random read or write. Thus this routine is of limited usefulness under CP/M. MS-DOS provides much more information however, and use of perror() after MS-DOS file handling calls will certainly give useful diagnostics.

SEE ALSO

open, close, read, write

PRINTF, VPRINTF

SYNOPSIS

```
#include <stdio.h>

int printf(char * fmt, ...)
int vprintf(char * fmt, va_list va_arg)
```

DESCRIPTION

Printf() is a formatted output routine, operating on stdout. There are corresponding routines operating on a given stream (fprintf()) or into a string buffer (sprintf()). Printf() is passed a format string, followed by a list of zero or more arguments. In the format string are conversion specifications, each of which is used to print out one of the argument list values. Each conversion specification is of the form %m.nc where the percent symbol % introduces a conversion, followed by an optional width specification m. n is an optional precision specification (introduced by the dot) and c is a letter specifying the type of the conversion. A minus sign ('-') preceding m indicates left rather than right adjustment of the converted value in the field. Where the field width is larger than required for the conversion, blank padding is performed at the left or right as specified. Where right adjustment of a numeric conversion is specified, and the first digit of m is 0, then padding will be performed with zeroes rather than blanks.

If the character * is used in place of a decimal constant, e.g. in the format %*d, then one integer argument will be taken from the list to provide that value. The types of conversion are:

- f Floating point - m is the total width and n is the number of digits after the decimal point. If n is omitted it defaults to 6.
- e Print the corresponding argument in scientific notation. Otherwise similar to f.
- g Use e or f format, whichever gives maximum precision in minimum width.
- o x X u d Integer conversion - in radices 8, 16, 10 and 10 respectively. The conversion is signed in the case of d, unsigned otherwise. The precision value is the total number of digits to print, and may be used to force leading zeroes. E.g. %8.4x will print at least 4 hex digits in an 8 wide field. Preceding the key letter with an l indicates that the value argument is a long integer or unsigned value. The letter X prints out hexadecimal numbers using the upper case letters A-F rather than a-f as would be printed when using x.

- s Print a string - the value argument is assumed to be a character pointer. At most n characters from the string will be printed, in a field m characters wide.
- c The argument is assumed to be a single character and is printed literally.

Any other characters used as conversion specifications will be printed. Thus %% will produce a single percent sign. Some examples:

```
printf("Total = %4d%%", 23)
    yields 'Total =23%'
printf("Size is %lx" , size)
    where size is a long, prints size
    as hexadecimal.
printf("Name = %.8s", "a1234567890")
    yields 'Name = a1234567'
printf("xx%d", 3, 4)
    yields 'xx 4'
```

Printf returns EOF on error, 0 otherwise. Vprintf() is similar to printf() but takes a variable argument list pointer rather than a list of arguments. See the description of va_start() for more information on variable argument lists.

SEE ALSO

fprintf, sprintf

PUTC

SYNOPSIS

```
#include <stdio.h>

int putc(int c, FILE * stream)
```

DESCRIPTION

Putc() is the macro version of fputc() and is defined in stdio.h. See fputc() for a description of its behaviour.

SEE ALSO

fputc, getc, fopen, fclose

PUTCHAR

SYNOPSIS

```
#include <stdio.h>

int putchar(int c)
```

DESCRIPTION

Putchar() is a putc() operation on stdout, defined in stdio.h.

SEE ALSO

putc, getc, freopen, fclose

PUTS

SYNOPSIS

```
#include <stdio.h>

int puts(char * s)
```

DESCRIPTION

Puts() writes the string *s* to the stdout stream, appending a newline. The null terminating the string is not copied. EOF is returned on error.

SEE ALSO

fputs, gets, freopen, fclose

PUTW

SYNOPSIS

```
#include <stdio.h>

int putw(int w, FILE * stream)
```

DESCRIPTION

Putw() copies the word *w* to the given stream. It returns *w*, except on error, in which case EOF is returned. Since this is a good integer, ferror() should be used to check for errors.

SEE ALSO

getw, fopen, fclose

QSORT

SYNOPSIS

```
#include <stdlib.h>

void      qsort(void * base, size_t nel,
               size_t width, int  (*func)())
```

DESCRIPTION

Qsort() is an implementation of the quicksort algorithm. It sorts an array of `nel` items, each of length `width` bytes, located contiguously in memory at `base`. `Func` is a pointer to a function used by `qsort()` to compare items. It calls `func` with pointers to two items to be compared. If the first item is considered to be greater than, equal to or less than the second then `func()` should return a value greater than zero, equal to zero or less than zero respectively.

```
static short  array[100];

#define SIZE  sizeof array/sizeof array[0]
a_func(p1, p2)
short * p1, * p2;
{
    return *p1 - *p2;
}

sort_em()
{
    qsort(array, SIZE,
          sizeof array[0], a_func);
}
```

This will sort the array into ascending values. Note the use of `sizeof` to make the code independent of the size of a `short`, or the number of elements in the array.

RAND

SYNOPSIS

```
#include <stdlib.h>

int rand(void)
```

DESCRIPTION

Rand() is a pseudo-random number generator. It returns an integer in the range 0 to 32767, which changes in a pseudo-random fashion on each call.

SEE ALSO

srand

READ

SYNOPSIS

```
#include <unistd.h>

int read(int fd, void * buf, size_t cnt)
```

DESCRIPTION

Read() will read from the file associated with fd up to cnt bytes into a buffer located at buf. It returns the number of bytes actually read. A zero return indicates end-of-file. A negative return indicates error. Fd should have been obtained from a previous call to open(). It is possible for read() to return less bytes than requested, e.g. when reading from the console, in which case read() will read one line of input.

SEE ALSO

open, close, write

REALLOC

SYNOPSIS

```
void * realloc(void * ptr, size_t cnt)
```

DESCRIPTION

Realloc() frees the block of memory at ptr, which should have been obtained by a previous call to malloc(), calloc() or realloc(), then attempts to allocate cnt bytes of dynamic memory, and if successful copies the contents of the block of memory located at ptr into the new block. At most, realloc() will copy the number of bytes which were in the old block, but if the new block is smaller, will only copy cnt bytes. If the block could not be allocated, 0 is returned.

SEE ALSO

malloc, calloc, realloc

REMOVE

SYNOPSIS

```
#include <stdio.h>

int remove(char * s)
```

DESCRIPTION

Remove() will attempt to remove the file named by the argument s from the directory. A return value of -1 indicates that the attempt failed.

SEE ALSO

unlink

RENAME

SYNOPSIS

```
#include <stdio.h>

int rename(char * name1, char * name2)
```

DESCRIPTION

The file named by name1 will be renamed to name2. -1 will be returned if the rename was not successful. Note that renames across user numbers or drives are not permitted.

SEE ALSO

open, close, unlink

REWIND

SYNOPSIS

```
#include <stdio.h>

int rewind(FILE * stream)
```

DESCRIPTION

This function will attempt to re-position the read/write pointer of the nominated stream to the beginning of the file. A return value of -1 indicates that the attempt was not successful, perhaps because the stream is associated with a non-random access file such as a character device.

SEE ALSO

fseek, ftell

SBRK

SYNOPSIS

```
char *    sbrk(int incr)
```

DESCRIPTION

Sbrk() increments the current highest memory location allocated to the program by incr bytes. It returns a pointer to the previous highest location. Thus sbrk(0) returns a pointer to the current highest location, without altering its value. If there is insufficient memory to satisfy the request, -1 is returned.

SEE ALSO

brk, malloc, calloc, realloc, free

SCANF

SYNOPSIS

```
#include <stdio.h>

int scanf(char * fmt, ...)
int vscanf(char *, va_list ap);
```

DESCRIPTION

Scanf() performs formatted input ("de-editing") from the stdin stream. Similar functions are available for streams in general, and for strings. The function vscanf() is similar, but takes a pointer to an argument list rather than a series of additional arguments. This pointer should have been initialized with va_start(). The input conversions are performed according to the fmt string; in general a character in the format string must match a character in the input; however a space character in the format string will match zero or more "white space" characters in the input, i.e. spaces, tabs or newlines. A conversion specification takes the form of the character %, optionally followed by an assignment suppression character ('*'), optionally followed by a numerical maximum field width, followed by a conversion specification character. Each conversion specification, unless it incorporates the assignment suppression character, will assign a value to the variable pointed at by the next argument. Thus if there are two conversion specifications in the fmt string, there should be two additional pointer arguments. The conversion characters are as follows:

o x d

Skip white space, then convert a number in base 8, 16 or 10 radix respectively. If a field width was

supplied, take at most that many characters from the input. A leading minus sign will be recognized.

- f Skip white space, then convert a floating number in either conventional or scientific notation. The field width applies as above.
- s Skip white space, then copy a maximal length sequence of non-white-space characters. The pointer argument must be a pointer to char. The field width will limit the number of characters copied. The resultant string will be null-terminated.
- c Copy the next character from the input. The pointer argument is assumed to be a pointer to char. If a field width is specified, then copy that many characters. This differs from the s format in that white space does not terminate the character sequence.

The conversion characters o, x, u, d and f may be preceded by an l to indicate that the corresponding pointer argument is a pointer to long or double as appropriate. A preceding h will indicate that the pointer argument is a pointer to short rather than int.

Scanf() returns the number of successful conversions; EOF is returned if end-of-file was seen before any conversions were performed. Some examples are:

```
scanf("%d %s", &a, &s)
    with input "12s"
will assign 12 to a, and "s" to s.

scanf("%4cd %lf", &c, &f)
    with input " abcd -3.5"
will assign " abcd" to c, and -3.5 to f.
```

SEE ALSO

fscanf, sscanf, printf, va_arg

SEGREAD

SYNOPSIS

```
#include <dos.h>
int segread(struct SREGS * segregs)
```

DESCRIPTION

Segread() copies the values of the segment registers into the structure pointed to by segregs.

SEE ALSO

int86, int86x, intdos, intdosx

SETJMP

SYNOPSIS

```
#include <setjmp.h>
int setjmp(jmp_buf buf)
```

DESCRIPTION

Setjmp() is used with longjmp() for non-local gotos. See longjmp() for further information.

SEE ALSO

longjmp

SETUID (CP/M only)

SYNOPSIS

```
#include <sys.h>

void setuid(int uid)
```

DESCRIPTION

Setuid() will set the current user number to uid. Uid should be a number in the range 0-15.

SEE ALSO

getuid

SETVBUF, SETBUF

SYNOPSIS

```
#include <stdio.h>

int setvbuf(FILE * stream, char * buf,
            int mode, size_t size);
void      setbuf(FILE * stream, char * buf)
```

DESCRIPTION

The setvbuf() function allows the buffering behaviour of a STDIO stream to be altered. It supersedes the function setbuf() which is retained for backwards compatibility. The arguments to setvbuf() are as follows: stream designates the STDIO stream to be affected; buf is a pointer to a buffer which will be used for all subsequent I/O operations on this stream. If buf is null, then the routine will allocate a buffer from the heap if necessary, of size BUFSIZ as defined in <stdio.h>. mode may take the values _IONBF, to turn buffering off completely, _IOFBF, for full buffering, or _IOLBF for line buffering. Full buffering means that the associated buffer will only be flushed when full, while line buffering means that the buffer will be flushed at the end of each line or when input is requested from another STDIO stream. size is the size of the buffer supplied. For example:

```
setvbuf(stdout, my_buf, _IOLBF, sizeof my_buf);
```

If a buffer is supplied by the caller, that buffer will remain associated with that stream even over fclose(), fopen() calls until another setvbuf() changes it.

SEE ALSO

fopen, freopen, fclose

SET_VECTOR

SYNOPSIS

```
#include <intrpt.h>
typedef interrupt void (*isr)();
isr set_vector(isr * vector, isr func);
```

DESCRIPTION

This routine allows an interrupt vector to be initialized. The first argument should be the address of the interrupt vector (not the vector number but the actual address) cast to a pointer to isr, which is a typedef'd pointer to an interrupt function. The second argument should be the function which you want the interrupt vector to point to. This must be declared using the interrupt type qualifier. The return value of set_vector() is the previous contents of the vector.

SEE ALSO

di(), ei()

SIGNAL

SYNOPSIS

```
#include <signal.h>
void (* signal)(int sig, void (*func)());
```

DESCRIPTION

Signal() provides a mechanism for catching control-C's (ctrl-BREAK for MS-DOS) typed on the console during I/O. Under CP/M the console is polled whenever an I/O call is performed, while for MS-DOS the polling depends on the setting of the BREAK command. If a control-C is detected certain action will be performed. The default action is to exit summarily; this may be modified with signal(). The sig argument to signal may at the present time be only SIGINT, signifying an interrupt condition. The func argument may be one of SIG_DFL, representing the default action, SIG_IGN, to ignore control-C's completely, or the address of a function which will be called with one argument, the number of the signal caught, when a control-C is seen. As the only signal supported is SIGINT, this will always be the value of the argument to the called function.

SEE ALSO

exit

SIN

SYNOPSIS

```
#include <math.h>

double    sin(double f);
```

DESCRIPTION

This function returns the sine function of its argument.

SEE ALSO

cos, tan, asin, acos, atan

SPAWN, SPAWNV, SPAWNVE

SYNOPSIS

```
int spawnl(char * n, char * argv0, ...);
int spawnv(char * n, char ** v)
int spawnve(char * n, char ** v, char ** e)
```

DESCRIPTION

These functions will load and execute a sub-program, named by the argument n. The calling conventions are similar to the functions execl() and execv(), the difference being that the spawn functions return to the calling program after termination of the sub-program, while the exec functions return only if the program could not be executed. Spawnve() takes an environment list in the same format as the argument list which will be supplied to the executed program as its environment.

SEE ALSO

execl, execv

SPRINTF

SYNOPSIS

```
#include <stdio.h>

int sprintf(char * buf, char * fmt, ...);
int vsprintf(char * buf, char * fmt, va_list ap);
```

DESCRIPTION

Sprintf() operates in a similar fashion to printf(), except that instead of placing the converted output on the stdout stream, the characters are placed in the buffer at buf. The resultant string will be null-terminated, and the number of characters in the buffer will be returned. Vsprintf takes an argument pointer rather than a list of arguments.

SEE ALSO

printf, fprintf, sscanf

SQRT

SYNOPSIS

```
#include <math.h>

double    sqrt(double f)
```

DESCRIPTION

Sqrt() implements a square root function using Newton's approximation.

SEE ALSO

exp

SSCANF

SYNOPSIS

```
#include <stdio.h>

int sscanf(char * buf, char * fmt, ...);
int vsscanf(char * buf, char * fmt, va_list ap);
```

DESCRIPTION

Sscanf() operates in a similar manner to scanf(), except that instead of the conversions being taken from stdin, they are taken from the string at buf.

SEE ALSO

scanf, fscanf, sprintf

SRAND

SYNOPSIS

```
#include <stdlib.h>

void srand(int seed)
```

DESCRIPTION

Srand() initializes the random number generator accessed by rand() with the given seed. This provides a mechanism for varying the starting point of the pseudo-random sequence yielded by rand(). On the z80, a good place to get a truly random seed is from the refresh register. Otherwise timing a response from the console will do.

SEE ALSO

rand

STAT

SYNOPSIS

```
#include <stat.h>

int stat(char * name, struct stat * statbuf)
```

DESCRIPTION

This routine returns information about the file by name. The information returned is operating system dependent, but may include file attributes (e.g. read only), file size in bytes, and file modification and/or access times. The argument name should be the name of the file, and may include path names under DOS, user numbers under CP/M, etc. The argument statbuf should be the address of a structure as defined in stat.h which will be filled in with the information about the file. The structure of struct stat is as follows:

```
struct stat
{
    short    st_mode; /* flags */
    long     st_atime; /* access time */
    long     st_mtime; /* modification time */
    long     st_size; /* file size */
};
```

The access and modification times (under DOS these are both set to the modification time) are in seconds since 00:00:00 Jan 1 1970. The function ctime() may be used to convert this to a readable

value. The file size is self explanatory. The flag bits are as follows:

Flag	Meaning
S_IFMT	mask for file type
S_IFDIR	file is a directory
S_IFREG	file is a regular file
S_IRREAD	file is readable
S_IWRITE	file is writeable
S_IXEXEC	file is executable
S_HIDDEN	file is hidden
S_SYSTEM	file is marked system
S_ARCHIVE	file has been written to

Stat returns 0 on success, -1 on failure, e.g. if the file could not be found.

SEE ALSO

ctime, creat, chmod

STRCAT, STRCMP, STRCPY, STRLEN et. al.

SYNOPSIS

```
#include <string.h>

char *   strcat(char * s1, char * s2);
int strcmp(char * s1, char * s2);
char *   strcpy(char * s1, char * s2);
int strlen(char * s);
char *   strncat(char * s1, char * s2, size_t n);
int strncmp(char * s1, char * s2, size_t n);
char *   strncpy(char * s1, char * s2, size_t n);
```

DESCRIPTION

These functions provide operations on null-terminated strings. Strcat() appends the string s2 to the end of the string s1. The string at s1 will be null terminated. Needless to say the buffer at s1 must be big enough. Strcmp() compares the two strings and returns a number greater than 0, 0 or a number less than 0 according to whether s1 is greater than, equal to or less than s2. The comparison is via the ascii collating order, with the first character the most significant. Strcpy() copies s2 into the buffer at s1, null terminating it. Strlen() returns the length of s1, not including the terminating null. Strncat(), strncmp() and strncpy() will concatenate, compare and copy s2 and s1 in the same manner as their similarly named counterparts above, but involving at most n characters. For strncpy(), the resulting string may not be null terminated.

STRCHR, STRRCHR

SYNOPSIS

```
#include <string.h>

char *   strchr(char * s, int c)
char *   strrchr(char * s, int c)
```

DESCRIPTION

These functions locate an instance of the character `c` in the string `s`. In the case of `strchr()` a pointer will be returned to the first instance of the character found by searching from the beginning of the string, while `strrchr()` searches backwards from the end of the string. A null pointer is returned if the character does not exist in the string.

SEE ALSO

SYSTEM

SYNOPSIS

```
#include <sys.h>

int system(char * s)
```

DESCRIPTION

When executed under MS-DOS `system()` will pass the argument string to the command processor, located via the environment string `COMSPEC`, for execution. The exit status of the command processor will be returned from the call to `system()`. For example, to set the baud rate on the serial port on an MS-DOS machine:

```
system("MODE COM1:96,N,8,1,P");
```

This function will not work on CP/M-86 since it does not have an invokable command interpreter. Under Concurrent CP/M the CLI system call is used.

SEE ALSO

`spawnl`, `spawnv`

TAN

SYNOPSIS

```
#include <math.h>

double    tan(double f);
```

DESCRIPTION

This is the tangent function.

SEE ALSO

sin, cos, asin, acos, atan

TIME

SYNOPSIS

```
#include <time.h>

time_t    time(time_t * t)
```

DESCRIPTION

This function returns the current time in seconds since 00:00:00 on Jan 1, 1970. If the argument *t* is non-null, the same value is stored into the object pointed to by *t*. The accuracy of this function is naturally dependent on the operating system having the correct time. This function does not work under CP/M-86 or CP/M 2.2 but does work under Concurrent-CP/M and CP/M+.

SEE ALSO

ctime, gmtime, localtime, asctime

TOUPPER, TOLOWER, TOASCII

SYNOPSIS

```
#include <ctype.h>
char toupper(int c);
char tolower(int c);
char toascii(int c);
char      c;
```

DESCRIPTION

Toupper() converts its lower case alphabetic argument to upper case, tolower() performs the reverse conversion, and toascii() returns a result that is guaranteed in the range 0-0177. Toupper() and tolower return their arguments if it is not an alphabetic character.

SEE ALSO

islower, isupper, isascii et. al.

UNGETC

SYNOPSIS

```
#include <stdio.h>

int ungetc(int c, FILE * stream)
```

DESCRIPTION

Ungetc() will attempt to push back the character c onto the named stream, such that a subsequent getc() operation will return the character. At most one level of pushback will be allowed, and if the stream is not buffered, even this may not be possible. EOF is returned if the ungetc() could not be performed.

SEE ALSO

getc

UNLINK

SYNOPSIS

```
int unlink(char * name)
```

DESCRIPTION

Unlink() will remove (delete) the named file, that is erase the file from its directory. See open() for a description of the file name construction. Zero will be returned if successful, -1 if the file did not exist or it could not be removed.

SEE ALSO

open, close, rename, remove

VA_START, VA_ARG, VA_END

SYNOPSIS

```
#include <stdarg.h>

void    va_start(va_list ap, parmN);
type    va_arg(ap, type);
void    va_end(va_list ap);
```

DESCRIPTION

These macros are provided to give access in a portable way to parameters to a function represented in a prototype by the ellipsis symbol (...), where type and number of arguments supplied to the function are not known at compile time. The rightmost parameter to the function (shown as parmN) plays an important role in these macros, as it is the starting point for access to further parameters. In a function taking variable numbers of arguments, a variable of type va_list should be declared, then the macro va_start invoked with that variable and the name of parmN. This will initialize the variable to allow subsequent calls of the macro va_arg to access successive parameters. Each call to va_arg requires two arguments; the variable previously defined and a type name which is the type that the next parameter is expected to be. Note that any arguments thus accessed will have been widened by the default conventions to int, unsigned int or double. For example if a character argument has been passed, it should be accessed by va_arg(ap, int) since the char will have been widened to int. An example is given below of a function taking one integer parameter, followed by a number of other parameters. In this example the function expects the subsequent parameters to be pointers to char, but note that the compiler is not aware of this, and it is the programmers responsibility to

ensure that correct arguments are supplied.

```
#include <stdarg.h>
prf(int n, ...)
{
    va_list      ap;

    va_start(ap, n);
    while(n--)
        puts(va_arg(ap, char *));
    va_end(ap);
}
```

WRITE

SYNOPSIS

```
#include <unixio.h>

int write(int fd, void * buf, size_t cnt)
```

DESCRIPTION

Write() will write from the buffer at buf up to cnt bytes to the file associated with the file descriptor fd. The number of bytes actually written will be returned. EOF or a value less than cnt will be returned on error. In any case, any return value not equal to cnt should be treated as an error (cf. read()).

SEE ALSO

open, close, read

INDEX

- 8086 70
- absolute value 116
- acos() 103
- address
 - link 73
 - load 73
- ANSI Standard 13,17
- array
 - dimension 16
 - index 16
 - size 16
- ASCII 26
- asctime() 104
- asin() 103
- assembler
 - in C programs 20
- assert() 105
- atan() 103
- atan2() 103
- Atari ST 3
- atexit() 104
- atof() 105
- atoi() 105
- atol() 105
- bdos() 106
- binary output file 8
- bios() 106
- braces
 - omission of 16
- byte ordering 21
- C command 3,7
- C reference books 2
- calloc() 107
- carriage return 26,120
- ceil() 116
- cgets() 107
- change file attributes 108
- character testing 134
- chdir() 108
- checksum 79,80
- chmod() 108
- class name 70
- close() 109
- clreof() 109
- clrerr() 109
- command line 7,12,77,160
- Compatibility 25
- compiler driver 3
- Compiler Structure 5
- console I/O 107,129,135
- control-C handling 155
- converting ascii to binary
 - 105
- cos() 110
- cosh() 110
- CP/M 26,106,131,155
- CP/M-80 compiler 2
- cputs() 107
- creat() 111
- creating a library 76
- CREF 83
- cross compilers 3
- cross reference 36,83
- ctime() 111
- ctrl-Z 26,120
- date and time 104
- debugging 10,105
- device name as file 23
- device
 - as file 141
- di() 113
- directory 108,139
 - current 130
- disk changing on floppy disk
 - systems 3
- disk space 33
- div() 112
- div_t 112
- dup() 113
- editor 3,33
- ei() 113
- entering long commands 7
- Enumerated Types 16
- environment 130
- error message:
 - Can't generate code 33
 - Error closing file 33
 - No room 33
 - out of memory 33
 - Undefined symbol 33
 - Write error 33
- error messages 23,33
 - format of 23
 - LIBR 78
 - list of 23,85
 - redirection of 23
- execl() 114
- executing a sub-program 156
- executing system commands
 - 160
- execution profiling 10
- execv() 114
- exit() 114
- exp() 115
- exponential functions 115
- Extern Declarations 30
- fabs() 116
- fclose() 116
- feof() 117
- ferror() 117

fflush() 117
fgetc() 118
fgets() 118
FILE 99
file descriptor 27,111,119
 duplicating 113
FILE pointer 27
file
 access time 158
 attributes 108,158
 closing a 109,116
 creating a 111
 deleting a 148,163
 flushing buffer to a 117
 information about a 158
 modification time 158
 opening a 119,123,140
 random access to a 125,137
 reading from a 122,147
 renaming a 149
 size 158
 type of a 135
 writing to a 126,164
fileno() 119
floating point
 library 27
 options required to use 4
floor() 116
fopen() 26,119
formatted input 124,150
formatted printing 121,142
fprintf() 121
fputc() 121
fputs() 122
fread() 122
free() 123
freopen() 123
frexp() 124
fscanf() 124
fseek() 125
ftell() 126
function parameter 163
function
 declaration 17
 large 33
 parameters 17
 prototype 17,163
fwrite() 126
getc() 128
getch() 129
getchar() 25,129
getche() 129
getcwd() 130
getenv() 130
gets() 131
getuid() 131
getw() 132
global variables 30
gmtime() 132
hex output file 8
I/O redirection 127
I/O
 ASCII 26
 Binary 26
 standard 25,99
in-line assembler 20
Initialization Syntax 16
initializer 16
inp() 133
int86() 133
int86x() 133
intdos() 133
intdosx() 133
Intel hex format 79
interrupt 133
interrupt handling 155
interrupt vector
 setting 154
isalnul() 134
isalnum() 134
isalpha() 134
isascii() 134
isatty() 135
iscntrl() 134
isdigit() 134
isgraph() 134
islower() 134
ispunct() 134
isspace() 134
isupper() 134
kbhit() 135
Kernighan and Ritchie 2,16
large model 70,80
ldexp() 124
ldiv() 112
LIBR 75
librarian 75
libraries 7
library manager 75
library ordering 77
library
 creating 76
 floating point 33
 ordering 33
line number symbols 10
linker 69
LINT 16
localtime 132
log() 115
log10() 115
logarithmic functions 115
long 30
longjmp() 136
lseek() 137

Machine Dependencies 21	pack pragma 20
macros	parameter
predefined 21	type 17
malloc() 137	perror() 141
Member Names 13,29	portable code 21
memcmp() 138	pow() 115
memcpy() 138	pragma 20,21
memory	printf 33
allocation	printf() 142
107,148,150	program execution 156
releasing 123	psect
memory model 31	linking 69
memory requirements 2	local 70
memset() 138	putc() 144
mkdir() 139	putch() 129
Motorola hex format 80	putchar() 144
MS-DOS 108,133,139,155	puts() 145
msdos() 139	putw() 145
msdoscx() 139	qsort() 146
multiple definitions 30	rand() 147
newline 26	random numbers 147,158
Newton's approximation 157	read() 147
non-local goto 136,152	realloc() 148
objtohex 72,79	relocation 69
open() 140	remove() 148
Operating Details 7	rename() 149
options 7	return 26,120
-1 10	rewind() 149
-11 11	rmdir() 139
-2 11	S format hex 80
-6301 11	sbrk() 150
-A 8	scanf 33
-B 9	scanf() 150
-C 7	segment registers 152
-CPM 7	segread() 152
-CR 7	self relocation 72
-D 8	setbuf() 153
-E 10	setjmp() 152
-F 8	setuid() 153
-G 10	setvbuf() 153
-H 10	set_vector() 154
-I 8	shift
-LF 4,27,33	signed 14
-M 8	unsigned 14
-O 8	short 21,30
-Ooutfile 8	signal() 155
-P 10	sin() 155
-R 8	sinh() 110
-S 7	size of data types 21
-U 8	sorting 146
-V 3,8	source file naming 3
-W 10	source level debugging 10
-X 8	spawnl() 156
-Z 10	spawnv() 156
linker 71	spawnve() 156
outp() 133	Specific Features 13
pack 21	sprintf() 156
	sqrt() 157

square root function 157	Unsigned Types 14
srand() 158	user ID 131,153
sscanf() 157	V3.09 2
Standard Libraries 25	va_arg 18,163
Standard Library Functions 99	va_end 163
stat() 158	va_start 163
stdarg.h 18,163	void 18
STDIO 25,27,99	Pointer to 18
storage allocator 18	warning level 10
strcat() 159	warning messages
strchr() 160	suppressing 10
strcmp() 159	wild card expansion 127
strcpy() 159	Wordstar 33
stream 27,99	write() 164
string manipulation 159	Z180 MMU 68
strlen() 159	Bank Base Register 68
strncat() 159	BBR 68
strncmp() 159	CBAR 68
strncpy() 159	CBR 68
strrchr() 160	Common Base Register 68
Structure Operations 15	Common/Bank Area Register 68
structure packing 20	Z180 Registers 68
structures	ASCI 68
as function argument 15	Bank Base Register 68
assignment 15	BBR 68
functions returning 15	BCR0H 68
Stylistic Considerations 29	BCR0L 68
symbol file 10,80	BCR1H 68
symbols	BCR1L 68
cross reference 83	CBAR 68
global 70,76	CBR 68
System Requirements 2	CNTLA0 68
system() 160	CNTLA1 68
tan() 161	CNTLB0 68
tanh() 110	CNTLB1 68
temporary files 4	CNTR 68
text editor 3	Common Base Register 68
time and date 104	Common/Bank Area Register 68
time() 161	CSI/0 68
toascii() 162	DAR0B 68
tolower() 162	DAR0H 68
toupper() 162	DAR0L 68
trigonometric functions	DCNTL 68
103,110	DMA 68
hyperbolic 110	DMA Byte Count 68
Type Checking 13	DMA Destination Address 68
type qualifiers 19	DMA Mode Register 68
typecast 13	DMA Source Address 68
typecasting 18	DMA Status Register 68
ungetc() 162	DMODE 68
ungetch() 129	DSTAT 68
Unix	FRC 68
C compiler 16	Free Running Counter 68
V7 25	IAR1H 68
unlink() 163	IAR1L 68
unsigned 23	
unsigned char 21	

ICR 68	INC 55
IL 68	IND 55
Interrupt Vector Register 68	INDR 56
ITC 68	INI 56
MAR1B 68	INIR 56
MAR1H 68	JP 56
MAR1L 68	JR 56
MMU 68	LD 56,57,58,59
OMCR 68	LDD 59
RCR 68	LDDR 59
Refresh Control Register 68	LDI 59
RLDR0H 68	LDIR 59
RLDR0L 68	MLT 59
RLDR1H 68	NEG 59
RLDR1L 68	NOP 59
SAR0B 68	OR 59,60
SAR0H 68	OTDM 61
SAR0L 68	OTDMR 60
STAT0 68	OTDR 60
STAT1 68	OTIM 61
TCR 68	OTIMR 60
TDR0 68	OTIR 60
TDR1 68	OUT 60
TIMER 0 68	OUT0 60
TIMER 1 68	OUTD 60
TMDR0H 68	OUTI 61
TMDR0L 68	POP 61
TMDR1H 68	PUSH 61
TMDR1L 68	RES 61,62,63
TRDR 68	RET 63
TSR0 68	RETI 63
TSR1 68	RETN 63
Z80 Instructions 50	RL 63
ADC 51	RLA 63
ADD 51	RLC 63
AND 51	RLCA 63
BIT 52,53	RLD 63
CALL 53	RR 63,64
CCF 53	RRA 64
CP 53,54	RRC 64
CPD 54	RRCA 64
CPDR 54	RRD 64
CPI 54	RST 64
CPIR 54	SBC 64
CPL 54	SCF 64
DAA 54	SET 64,65,66
DEC 54	SLA 66
DI 54	SLP 66
DJNZ 54	SRA 66
EI 54	SRL 67
EX 55	SUB 67
EXX 55	TST 67
HALT 55	TSTIO 67
IM 55	XOR 67
IN 55	Z80 pseudo ops
IN0 55	IRP 47
	IRPC 47
	ZAS directives 48

title 48

ZAS options

- J 35
- L 36
- N 35
- O 36
- U 35
- W 36

ZAS psect flags

- ABS 43
- GLOBAL 43
- LOCAL 43
- OVRD 43
- PURE 43

ZAS pseudo ops 40

- COND 42
- DB 40
- DEFB 40
- DEFF 41
- DEFL 41
- DEFM 42
- DEFS 41
- DEFW 41
- ELSE 42
- END 42
- ENDC 42
- ENDM 44
- EQU 41
- GLOBAL 43
- IF 42
- LOCAL 45
- MACRO 44
- ORG 44
- PSECT 43
- REPT 46

ZAS

- 64180 35
- arithmetic overflow 35
- binary constants 37
- character constants 38
- condition codes 48
- conditional assembly 42
- constants 37
- directives 48
- extended condition codes 48
- external symbols 35
- floating point constants 38
- hexadecimal constants 37
- jump optimization 35
- jumps 35
- labels 36
- listing 36
- macros 44
- mnemonics 35
- octal constants 37
- opcode constants 38
- operators 38
- program sections 39
- PSECT 43
- psects 39,43
- pseudo ops 40
- relocatable object code 39
- temporary labels 37
- Z180 35
- Zilog 35
- _exit() 115
- _getargs() 127